

UNIVERSITY OF MUMBAI

PRACTICAL REPORT
M.Sc. I.T (PART 2) 2024-2025


PAPER 1: Advance Artificial Intelligence

PAPER 2: Machine Learning

**SUBMITTED BY:
(YASIN KHAN)**

**(PROF. SIMRAN A. SHAIKH)
{TEACHER INCHARGE}**

**(PROF. SRAVANTI GAJJELLI)
{TEACHER INCHARGE}**



**Padmashri Annasaheb Jadhav Bhartiya Samaj Unnati Mandal's
B.N.N COLLEGE**

(Arts, Science, Commerce & Self-Funded Course)

(Affiliated to University of Mumbai)

**DEPARTMENT OF INFORMATION TECHNOLOGY BHIWANDI,
MAHARASHTRA-421302**



**Padmashri Annasaheb Jadhav Bhartiya Samaj Unnati Mandal's
B.N.N College of Arts, Science and Commerce Bhiwandi
(Self-Funded Course)**

(Department of Information Technology)

CERTIFICATE

This is to certify that Mr./Miss. _____

Class: **MSc-IT** Exam Seat No. 13132 has satisfactorily completed practical in. **Advance Artificial Intelligence**. As laid down in the regulation of University of Mumbai for the purpose of **Semester- III Practical** Examination 2024-2025.

Date:

Place: **Bhiwandi**

In-Charge Professor

Signature of External Examiner

Signature of HOD

Signature of Principal



Padmashri Annasaheb Jadhav Bhartiya Samaj Unnati Mandal's
B.N.N College of Arts, Science and Commerce Bhiwandi
(Self-Funded Course)

(Department of Information Technology)

CERTIFICATE

This is to certify that Mr./Miss. _____

Class: **MSc-IT** Exam Seat No. 13123. has satisfactorily completed practical in. **Machine Learning**. As laid down in the regulation of University of Mumbai for the purpose of **Semester-III Practical** Examination 2024-2025.

Date:

Place: **Bhiwandi**

In-Charge Professor

Signature of External Examiner

Signature of HOD

Signature of Principal

Sr. No.	Title	Date	Sign
1	Implementing advanced deep learning algorithms such as convolutional neural networks (CNNs) or recurrent neural networks (RNNs) using Python libraries like TensorFlow or PyTorch.	07-09-2025	
2	Building a natural language processing (NLP) model for sentiment analysis or text classification.	10-09-2025	
3	Creating a chatbot using advanced techniques like transformer models.	20-09-2025	
4	Developing a recommendation system using collaborative filtering or deep learning approaches.	26-09-2025	
5	Implementing a computer vision project, such as object detection or image segmentation.	04-10-2025	
6	Training a generative adversarial network (GAN) for generating realistic images.	11-10-2025	
7	Applying reinforcement learning algorithms to solve complex decision-making problems.	19-10-2025	
8	Utilizing transfer learning to improve model performance on limited datasets.	26-10-2025	
9	Building a deep learning model for time series forecasting or anomaly detection.	02-11-2025	
10	Implementing a machine learning pipeline for automated feature engineering and model selection.	16-11-2025	
11	Using advanced optimization techniques like evolutionary algorithms or Bayesian optimization for hyperparameter tuning.	30-11-2025	
12	Deploying a machine learning model in a production environment using containerization and cloud services.	06-12-2025	
13	Use Python libraries such as GPT-2 or textgenrnn to train generative models on a corpus of text data and generate new text based on the patterns it has learned.	14-12-2025	
14	Experiment with neural networks like GANs (Generative Adversarial Networks) using Python libraries like TensorFlow or PyTorch to generate new images based on a dataset of images.	28-12-2025	

Sr. No.	Title	Date	Sign
1A	Load a CSV dataset. Handle missing values, inconsistent formatting, and outliers.	06-09-2025	
1B	Load a dataset, calculate descriptive summary statistics, create visualizations using different graphs, and identify potential features and target variables Note: Explore Univariate and Bivariate graphs (Matplotlib) and Seaborn for visualization.	14-09-2025	
1C	Create or Explore datasets to use all pre-processing routines like label encoding, scaling, and binarization.	20-09-2025	
2	Implement and demonstrate the FIND-S algorithm for finding the most specific hypothesis based on a given set of training data samples. Read the training data from a .CSV file and generate the final specific hypothesis. (Create your dataset)	28-09-2025	
3A	Simple Linear Regression Fit a linear regression model on a dataset. Interpret coefficients, make predictions, and evaluate performance using metrics like R-squared and MSE	05-10-2025	
3B	Extend linear regression to multiple features. Handle feature selection and potential multicollinearity.	12-10-2025	
3C	Implement regression variants like LASSO and Ridge on any generated dataset.	19-10-2025	
4A	Perform binary classification using logistic regression. Calculate accuracy, precision, recall, and understand the ROC curve.	25-10-2025	
4B	Implement and demonstrate k-nearest Neighbor algorithm. Read the training data from a .CSV file and build the model to classify a test sample. Print both correct and wrong predictions.	09-11-2025	
4C	Build a decision tree classifier or regressor. Control hyperparameters like tree depth to avoid overfitting. Visualize the tree.	15-11-2025	
4D	Implement a Support Vector Machine for any relevant dataset.	23-11-2025	
4E	Train a random forest ensemble. Experiment with the number of trees and feature sampling. Compare performance to a single decision tree.	29-11-2025	
4F	Implement a gradient boosting machine (e.g., XGBoost). Tune hyperparameters and explore feature importance.	06-12-2025	
5A	Implement and demonstrate the working of a Naive Bayesian classifier using a sample data set. Build the model to classify a test sample.	07-12-2025	
5B	Implement Hidden Markov Models using hmmlearn	13-12-2025	
6A	Implement Bayesian Linear Regression to explore prior and posterior distribution.	14-12-2025	
6B	Implement Gaussian Mixture Models for density estimation and unsupervised clustering	21-12-2025	
7A	Implement cross-validation techniques (k-fold, stratified, etc.) for robust model evaluation	28-12-2025	
7B	Systematically explore combinations of hyperparameters to optimize model performance.(use grid and randomized search)	04-12-2025	
8	Implement Bayesian Learning using inferences	10-12-2025	
9	Set up a generator network to produce samples and a discriminator network to distinguish between real and generated data. (Use a simple small dataset)	18-12-2025	
10	Develop an API to deploy your model and perform predictions	24-12-2025	

1. Convolutional Neural Network (CNN) for Image Classification

CNNs are widely used in computer vision tasks, such as image classification, object detection, and segmentation. They are particularly effective for processing grid-like data (e.g., images) due to their ability to capture spatial hierarchies.

We'll implement a CNN for image classification on the **MNIST** dataset, a collection of handwritten digits.

Step 1: Install Dependencies

Install the required libraries if you don't have them:

```
pip install tensorflow matplotlib numpy
```

Step 2: Load and Preprocess the Dataset

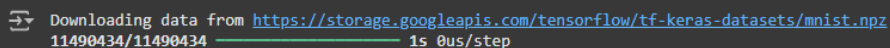
We'll use the **MNIST** dataset, which is available directly in TensorFlow.

```
import tensorflow as tf
from tensorflow.keras import layers, models
import numpy as np
import matplotlib.pyplot as plt

# Load MNIST dataset
(x_train, y_train), (x_test, y_test) = tf.keras.datasets.mnist.load_data()

# Normalize the images to [0, 1] and reshape them to (28, 28, 1)
x_train = x_train.astype('float32') / 255.0
x_test = x_test.astype('float32') / 255.0
x_train = np.expand_dims(x_train, axis=-1) # Add channel dimension
x_test = np.expand_dims(x_test, axis=-1) # Add channel dimension

# One-hot encode the labels
y_train = tf.keras.utils.to_categorical(y_train, 10)
y_test = tf.keras.utils.to_categorical(y_test, 10)
```



A terminal window showing the download of the MNIST dataset. The text indicates that data is being downloaded from a Google Cloud Storage link. The progress bar shows that approximately 15% of the data has been downloaded, with a speed of 0 bytes per step.

Step 3: Define the CNN Model

We will create a simple CNN architecture with 2 convolutional layers, followed by a fully connected (dense) layer.

```
def build_cnn_model():
    model = models.Sequential()

    # First convolutional layer
    model.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28,
1)))
    model.add(layers.MaxPooling2D((2, 2)))

    # Second convolutional layer
    model.add(layers.Conv2D(64, (3, 3), activation='relu'))
    model.add(layers.MaxPooling2D((2, 2)))

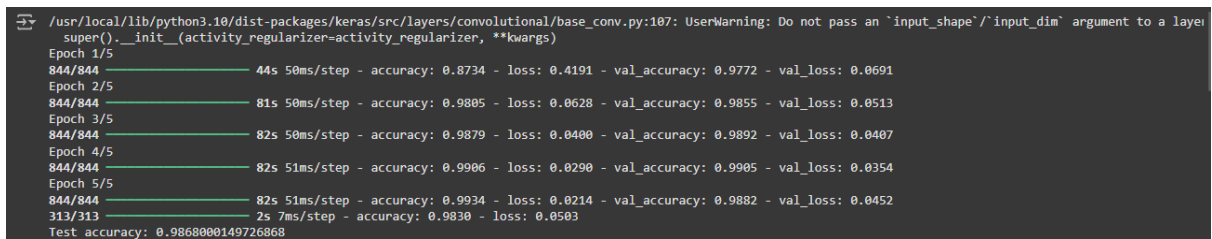
    # Flatten the output for the dense layer
    model.add(layers.Flatten())

    # Fully connected layer
    model.add(layers.Dense(64, activation='relu'))

    # Output layer
    model.add(layers.Dense(10, activation='softmax')) # 10 classes for MNIST
digits

    return model

# Build and compile the model
cnn_model = build_cnn_model()
cnn_model.compile(optimizer='adam',
                  loss='categorical_crossentropy',
                  metrics=['accuracy'])
```



```
/usr/local/lib/python3.10/dist-packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
Epoch 1/5
844/844 — 44s 50ms/step - accuracy: 0.8734 - loss: 0.4191 - val_accuracy: 0.9772 - val_loss: 0.0691
Epoch 2/5
844/844 — 81s 50ms/step - accuracy: 0.9805 - loss: 0.0628 - val_accuracy: 0.9855 - val_loss: 0.0513
Epoch 3/5
844/844 — 82s 50ms/step - accuracy: 0.9879 - loss: 0.0400 - val_accuracy: 0.9892 - val_loss: 0.0407
Epoch 4/5
844/844 — 82s 51ms/step - accuracy: 0.9906 - loss: 0.0290 - val_accuracy: 0.9905 - val_loss: 0.0354
Epoch 5/5
844/844 — 82s 51ms/step - accuracy: 0.9934 - loss: 0.0214 - val_accuracy: 0.9882 - val_loss: 0.0452
313/313 — 2s 7ms/step - accuracy: 0.9830 - loss: 0.0503
Test accuracy: 0.9868000149726868
```

Step 4: Train the Model

Now, we'll train the CNN on the MNIST dataset.

```
# Train the CNN model
cnn_model.fit(x_train, y_train, epochs=5, batch_size=64, validation_split=0.1)
```

Step 5: Evaluate the Model

After training, we'll evaluate the model on the test set.


```
# Evaluate the model on the test data
test_loss, test_acc = cnn_model.evaluate(x_test, y_test)
print(f'Test accuracy: {test_acc}')
```


Step 6: Visualize the Results



You can visualize the results by plotting the predictions made by the CNN.

```
# Make predictions
predictions = cnn_model.predict(x_test)

# Display the first 5 images and their predicted labels
for i in range(5):
    plt.imshow(x_test[i].reshape(28, 28), cmap='gray')
    plt.title(f"Predicted: {np.argmax(predictions[i])}, Actual:
{np.argmax(y_test[i])}")
    plt.show()
```





2. Recurrent Neural Network (RNN) for Sequence Prediction

RNNs are well-suited for sequential data like time-series, text, or audio. They maintain hidden states over time, making them effective for sequence modeling.

Let's implement a simple RNN using TensorFlow to predict the next word in a sequence.

Step 1: Install Dependencies

If you don't have the required libraries yet, install them:

```
pip install tensorflow numpy
```

Step 2: Prepare the Dataset

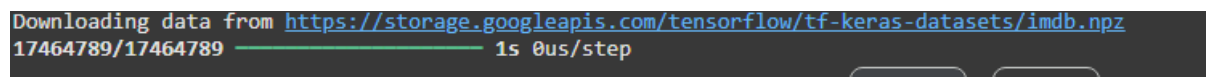
We'll use the **IMDB dataset** (a movie review dataset) for text classification. We'll build an RNN to predict the sentiment of a movie review (positive or negative).

```
# Load IMDB dataset for sentiment analysis
from tensorflow.keras.datasets import imdb
from tensorflow.keras.preprocessing.sequence import pad_sequences

# Set maximum number of words to consider and maximum sequence length
max_features = 10000 # Only consider the top 10,000 words
maxlen = 500 # Maximum length of the review

# Load and preprocess the dataset
(x_train, y_train), (x_test, y_test) = imdb.load_data(num_words=max_features)

# Pad sequences to ensure they have the same length
x_train = pad_sequences(x_train, maxlen=maxlen)
x_test = pad_sequences(x_test, maxlen=maxlen)
```



```
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/imdb.npz
17464789/17464789 ————— 1s 0us/step
```

Step 3: Define the RNN Model

We will define an RNN with **LSTM** layers. LSTM (Long Short-Term Memory) is an advanced type of RNN that solves the vanishing gradient problem, making it more effective for long sequences.

```
from tensorflow.keras import models, layers # Import necessary modules
```

```
def build_rnn_model():
    model = models.Sequential()
```

```

    # Embedding layer to learn word representations
    model.add(layers.Embedding(input_dim=max_features, output_dim=128,
                               input_length=maxlen))

    # LSTM layer
    model.add(layers.LSTM(128))

    # Output layer (sigmoid for binary classification)
    model.add(layers.Dense(1, activation='sigmoid'))

    return model

# Build and compile the RNN model
rnn_model = build_rnn_model()
rnn_model.compile(optimizer='adam', loss='binary_crossentropy',
                  metrics=['accuracy'])

```

Step 4: Train the Model

We will train the RNN model on the IMDB dataset.

```

# Train the RNN model
rnn_model.fit(x_train, y_train, epochs=5, batch_size=64,
              validation_data=(x_test, y_test))

```

```

Epoch 1/5
391/391 ————— 672s 2s/step - accuracy: 0.6965 - loss: 0.5553 - val_accuracy: 0.6059 - val_loss: 0.6345
Epoch 2/5
391/391 ————— 719s 2s/step - accuracy: 0.8072 - loss: 0.4287 - val_accuracy: 0.8678 - val_loss: 0.3175
Epoch 3/5
391/391 ————— 714s 2s/step - accuracy: 0.9075 - loss: 0.2451 - val_accuracy: 0.8703 - val_loss: 0.3213
Epoch 4/5
391/391 ————— 715s 2s/step - accuracy: 0.9438 - loss: 0.1599 - val_accuracy: 0.8734 - val_loss: 0.3277
Epoch 5/5
391/391 ————— 702s 2s/step - accuracy: 0.9385 - loss: 0.1654 - val_accuracy: 0.8710 - val_loss: 0.3555
<keras.src.callbacks.history.History at 0x7935dfa2f7f0>

```

Step 5: Evaluate the Model

After training, we can evaluate the model's performance on the test data.

```

# Evaluate the model on the test data
test_loss, test_acc = rnn_model.evaluate(x_test, y_test)
print(f'Test accuracy: {test_acc}')

```

```

782/782 ————— 228s 291ms/step - accuracy: 0.8685 - loss: 0.3606
Test accuracy: 0.8709999918937683

```

Step 6: Make Predictions

Once the model is trained, we can use it to make predictions on new reviews.

```
# Predict sentiment for the first review in the test set
predictions = rnn_model.predict(x_test[:5])

# Print predictions
for i, prediction in enumerate(predictions):
    sentiment = 'positive' if prediction >= 0.5 else 'negative'
    print(f"Review {i+1}: {sentiment} (probability: {prediction[0]:.2f})")
```

Conclusion

In this we've demonstrated how to implement advanced deep learning algorithms using **TensorFlow**:

1. **CNN**: We used a convolutional neural network to classify handwritten digits from the MNIST dataset. CNNs are powerful for image recognition tasks, and they work by learning spatial hierarchies through convolution and pooling operations.
2. **RNN**: We used a recurrent neural network (RNN) with LSTM units to classify sentiment from movie reviews in the IMDB dataset. RNNs are ideal for sequence data because they maintain hidden states across time steps, enabling them to capture dependencies in the data.

These models are just starting points. You can extend them with more complex architectures like **ResNets**, **Inception Networks**, or **Transformers** (for sequence-to-sequence tasks), and integrate them with other techniques like **transfer learning** or **data augmentation** for even better performance.

PRACTICAL NO. 2:

BUILDING A NLP MODEL FOR SENTIMENT ANALYSIS OR TEXT CLASSIFICATION

Here's a step-by-step approach to building a sentiment analysis model using real-world data from IMDb. We will:

1. Load the dataset (IMDb movie reviews).
2. Preprocess the text.
3. Convert the text to features using TF-IDF.

4. Train a sentiment analysis model.
5. Evaluate the model.

Explanation of the Code:

1. Libraries:

- **nltk**: Used for tokenizing and removing stop words.
- **sklearn**: Provides tools for vectorizing the text, splitting data into training and test sets, and training a classifier.
- **load_files**: Used to load a folder of text files containing reviews (positive and negative).
- **MultinomialNB**: Naive Bayes classifier that works well with word frequencies.

2. Dataset:

- We use the **IMDb dataset** that is commonly used for sentiment analysis, containing movie reviews classified as **positive (1)** and **negative (0)**.
- The dataset consists of text files organized in two directories: **pos** and **neg** (for positive and negative reviews).

3. Preprocessing:

- The text is first converted to lowercase to standardize it.
- Tokenization splits the text into words.
- Stopwords are removed using the NLTK library, which filters out common words like “the”, “is”, etc.
- Only alphabetic words are kept, removing any numbers or special characters.

4. Feature Extraction:

- The **TF-IDF Vectorizer** converts the cleaned text into numerical data that a machine learning model can process. It assigns importance to words based on their frequency in the document and across the entire dataset.

5. Model Training:

- We split the dataset into training and testing sets (80% for training, 20% for testing).
- We train the **Naive Bayes** model (**MultinomialNB**) using the training data.

6. Evaluation:

- We calculate the accuracy of the model using **accuracy_score**, which tells us how well the model classified the test data.
- The **confusion matrix** shows the breakdown of correct and incorrect predictions (True Positives, True Negatives, False Positives, False Negatives).

7. Accuracy: This is the percentage of correctly classified reviews (positive and negative).

8. Confusion Matrix: This shows the number of true positives (correctly classified positive reviews), true negatives (correctly classified negative reviews), false positives (negative

reviews incorrectly classified as positive), and false negatives (positive reviews incorrectly classified as negative).

Notes:

- **Data Source:** In this example, the `load_files` function is used to load the dataset from a directory. For real applications, you can download the IMDb dataset from Kaggle or another source and store it locally.
- **Model Improvements:** The Naive Bayes model can be further tuned, or you can try more advanced models like **Logistic Regression** or **Deep Learning models** (like **LSTM** or **BERT**) for better accuracy.

Steps for Making Predictions

1. **Define the prediction function:** This function will take a new review as input, preprocess the text, vectorize it, and use the model to predict the sentiment.
2. **Input a review:** Provide a review and get the sentiment prediction (positive or negative).

Explanation of the Code:

1. **predict_sentiment function:**
 - The function takes a review as input.
 - It preprocesses the text using the same `preprocess_text()` function as during training (to ensure the input is in the same format as the training data).
 - It then transforms the processed text into a vector using the trained **TF-IDF vectorizer** (`vectorizer.transform`).
 - The transformed vector is passed into the trained **Naive Bayes model** (`model.predict`) to predict the sentiment (either `1` for positive or `0` for negative).
 - The function returns the sentiment as either **Positive** or **Negative** based on the prediction.
2. **Making the Prediction:**
 - The new review, "I absolutely loved this movie! The plot was amazing and the acting was superb.", is passed to the `predict_sentiment()` function.
 - The function processes the review and outputs the sentiment prediction.

Important Notes:

- **Preprocessing Consistency:** It's crucial that the new review is preprocessed in the same way as the training data. The `preprocess_text` function does this by converting the text to lowercase, removing non-alphabetic tokens, and filtering out stop words.
- **TF-IDF Vectorization:** Since you trained the **TF-IDF vectorizer** on the training data, the new review must be transformed using the same vectorizer (`vectorizer.transform`). This ensures that the model recognizes the words in the same way it did during training.

Step 1: Import necessary libraries

```
!pip install datasets
import nltk
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.naive_bayes import MultinomialNB
from sklearn.metrics import accuracy_score, confusion_matrix
from sklearn.datasets import load_files
from datasets import load_dataset
```

Download NLTK resources (if not already installed)

```
nltk.download('stopwords')
nltk.download('punkt')
nltk.download('punkt_tab') # Download the 'punkt_tab' data
```

Step 2: Load the IMDb dataset (from sklearn)

IMDb dataset has positive and negative reviews stored in "pos" and "neg" folders

Load the dataset

```
from datasets import load_dataset # Import load_dataset from the datasets library
dataset = load_dataset('imdb')
print(dataset)
```

Step 3: Preprocessing the Text

This function preprocesses the text by converting it to lowercase and removing non-alphabetic words.

```
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize
```

```
stop_words = set(stopwords.words('english'))
```

```
def preprocess_text(text):
    text = text.lower() # Convert text to lowercase
    words = word_tokenize(text) # Tokenize the text
    words = [word for word in words if word.isalpha()] # Remove non-alphabetic tokens
    words = [word for word in words if word not in stop_words] # Remove stopwords
    return " ".join(words)
```

Preprocess all the reviews

```
processed_texts = [preprocess_text(text) for text in dataset['train']['text']]
```

Step 4: Feature Extraction using TF-IDF Vectorizer

```
vectorizer = TfidfVectorizer(max_features=1000) # Limit to 1000 features (words)
X = vectorizer.fit_transform(processed_texts) # Convert text data to feature vectors
y = dataset['train']['label'] # Sentiment labels (1 = Positive, 0 = Negative)
```

Step 5: Train-test split (80% train, 20% test)

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

Step 6: Train a model (Naive Bayes Classifier)

```
model = MultinomialNB()
model.fit(X_train, y_train)
```

Step 7: Make Predictions

```
y_pred = model.predict(X_test)
```

Step 8: Evaluate the Model

```
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy * 100:.2f}%")
```

Confusion Matrix

```
cm = confusion_matrix(y_test, y_pred)
print("Confusion Matrix:")
```

```
print(cm)

# Step 1: Define a function for prediction
def predict_sentiment(review_text):
    # Preprocess the text (same preprocessing steps used during training)
    processed_review = preprocess_text(review_text)

    # Convert the preprocessed text into features using the trained TF-IDF vectorizer
    review_vector = vectorizer.transform([processed_review])

    # Predict sentiment using the trained model
    prediction = model.predict(review_vector)

    # Return the sentiment prediction (1 = Positive, 0 = Negative)
    if prediction == 1:
        return "Positive"
    else:
        return "Negative"

# Step 2: Make a prediction with a new review
new_review = "I absolutely loved this movie! The plot was amazing and the acting was superb."

# Get the prediction for the new review
predicted_sentiment = predict_sentiment(new_review)
print(f"The sentiment of the review is: {predicted_sentiment}")
```

OUTPUT

Accuracy: 83.42%

Confusion Matrix:

```
[[2053 462]
```

```
 [ 367 2118]]
```

The sentiment of the review is: Positive

2. CODE FOR TEXT CLASSIFICATION

Explanation of the Code:

1. Dataset:

- We use the 20 Newsgroups dataset (`fetch_20newsgroups`) from `scikit-learn`, which contains 20 different categories of news articles. Each category represents a different type of news (e.g., sports, politics, technology, etc.).
- This dataset contains multiple categories, so it's a multi-class classification problem.

2. Text Preprocessing:

- The text preprocessing steps are similar to the previous example. We convert the text to lowercase, tokenize the text, remove non-alphabetic words, and remove stopwords to clean the data.

3. Feature Extraction:

- We use TF-IDF Vectorizer to transform the text into numerical feature vectors. The `max_features=1000` argument limits the features to the top 1000 words based on their importance.

4. Model:

- We use Naive Bayes (MultinomialNB) for classification, which is a popular choice for text classification tasks. It works well with discrete data (like word counts).

5. Training and Evaluation:

- The dataset is split into a training set (80%) and a test set (20%).
- We train the Naive Bayes model on the training data and then make predictions on the test data.
- We evaluate the model using accuracy and a confusion matrix to understand how well the model classifies text into the correct categories.

6. Making Predictions:

- The function `predict_category` allows you to input a new article (a string of text), preprocess it, transform it into a feature vector, and then use the trained model to predict its category.

```
# Step 1: Import necessary libraries
```

```

import nltk
from sklearn.datasets import fetch_20newsgroups
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.naive_bayes import MultinomialNB
from sklearn.metrics import accuracy_score, confusion_matrix

# Download NLTK resources (if not already installed)
nltk.download('stopwords')
nltk.download('punkt')
nltk.download('punkt_tab') # Download the punkt_tab resource

# Step 2: Load the 20 Newsgroups Dataset (multi-class text classification)
# Fetch the dataset (contains 20 categories of news articles)
newsgroups = fetch_20newsgroups(subset='all', remove=('headers',
'footers', 'quotes'))

# Step 3: Preprocessing the Text
# This function preprocesses the text by converting it to lowercase and
removing non-alphabetic words.
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize

stop_words = set(stopwords.words('english'))

def preprocess_text(text):
    text = text.lower() # Convert text to lowercase
    words = word_tokenize(text) # Tokenize the text
    words = [word for word in words if word.isalpha()] # Remove
non-alphabetic tokens
    words = [word for word in words if word not in stop_words] # Remove
stopwords
    return " ".join(words)

# Preprocess all the articles
processed_texts = [preprocess_text(text) for text in newsgroups.data]

```

```

# Step 4: Feature Extraction using TF-IDF Vectorizer
vectorizer = TfidfVectorizer(max_features=1000) # Limit to 1000 features
(words)
X = vectorizer.fit_transform(processed_texts) # Convert text data to
feature vectors
y = newsgroups.target # Categories (labels) for the newsgroups

# Step 5: Train-test split (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)

# Step 6: Train a model (Naive Bayes Classifier)
model = MultinomialNB()
model.fit(X_train, y_train)

# Step 7: Make Predictions
y_pred = model.predict(X_test)

# Step 8: Evaluate the Model
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy * 100:.2f}%")

# Confusion Matrix
cm = confusion_matrix(y_test, y_pred)
print("Confusion Matrix:")
print(cm)

# Step 9: Making Predictions for new articles
def predict_category(article_text):
    processed_article = preprocess_text(article_text)
    article_vector = vectorizer.transform([processed_article])
    prediction = model.predict(article_vector)
    return newsgroups.target_names[prediction[0]]

# Example of a new article
new_article = "The economy is showing signs of recovery. Experts suggest
that the job market will improve in the coming months."

```

```

predicted_category = predict_category(new_article)
print(f"The predicted category for the article is: {predicted_category}")

```

OUTPUT:

```

Accuracy: 54.30%
Confusion Matrix:
[[ 58   0   2   0   0   2   0   2  12   3   0   1   1  12   4  32   8   7
   6   1]
 [  3  95  23  12   2  13   5   2   9   3   2   1  14   8   7   2   0   1
   0   0]
 [  1  20 102  24   5  13   0   2  11   1   1   3   2   7   1   2   0   0
   0   0]
 [  0  14  20  91  22   5   7   5   3   1   0   0  10   2   1   1   1   0
   0   0]
 [  1  10  10  37  93   6   4   5  16   0   0   3   7   5   5   3   0   0
   0   0]
 [  0  35  15   7   3 140   1   3   4   1   0   1   1   0   2   1   1   0
   0   0]
 [  0   2   3  19  11   3 126   8   4   3   0   3   7   1   2   0   1   0
   0   0]
 [  6   2   1   3   3   4   5  98  34   4   2   2   7   6   8   2   8   0
   1   0]
 [  1   3   1   4   0   3   6  17  96  10   4   1   1   8   2   2   5   2
   2   0]
 [  4   3   1   0   1   1   2   6  21 114  36   1   3   6   3   4   0   5
   0   0]

```

```

[ 4 1 0 0 0 3 2 1 14 30 133 2 2 2 1 2 1 0
 0 0]
[ 3 7 1 0 5 1 2 2 5 0 5 123 14 10 5 6 5 3
 4 0]
[ 1 13 8 14 4 7 10 11 14 1 1 5 94 8 7 2 2 0
 0 0]
[ 6 9 1 0 2 1 1 3 13 1 6 1 5 128 2 8 3 1
 3 0]
[ 2 10 1 1 3 3 1 5 20 6 2 5 6 6 111 2 3 1
 1 0]
[ 8 2 0 0 2 1 1 0 8 3 1 0 2 5 1 158 2 5
 2 1]
[ 5 0 1 0 0 1 0 3 18 4 4 5 0 3 5 7 114 2
 15 1]
[ 15 0 2 0 1 1 1 0 14 3 3 3 1 9 2 9 5 107
 6 0]
[ 5 0 1 0 0 1 1 4 12 2 3 1 1 8 10 18 24 5
 63 0]
[ 22 1 0 0 0 1 0 0 14 3 3 1 1 11 3 57 8 5
 3 3]]
The predicted category for the article is: rec.autos

```

3. CREATING A CHATBOT USING ADVANCED TECHNIQUES LIKE TRANSFORMERS MODELS.

EXPLANATION :

Creating a chatbot using Transformer models involves using pre-trained models from libraries such as Hugging Face's Transformers. These models, such as GPT, BERT, or T5, have revolutionized the field of NLP due to their effectiveness in understanding and generating human language.

In this guide, we'll focus on creating a simple chatbot using a pre-trained transformer model, specifically DialoGPT, which is fine-tuned for conversation.

Steps to Create a Chatbot Using Hugging Face's **transformers** Library

1. Install necessary libraries:

- You need the **transformers** library from Hugging Face and **torch** (PyTorch) to load and run transformer models.

- `pip install transformers torch`

Import required libraries:

- We'll load a pre-trained model and tokenizer from Hugging Face's `DialoGPT` (a variant of GPT fine-tuned on conversation datasets).

Chatbot Setup:

- Load the pre-trained model.
- Tokenize input text, generate the response, and decode the output.
- Set up a loop for ongoing conversation.

Explanation of Code:

1. Libraries:
 - `GPT2LMHeadModel` and `GPT2Tokenizer` from the `transformers` library are used to load the pre-trained DialoGPT model and tokenizer.
 - `torch` is used for tensor operations (since Hugging Face's models are based on PyTorch).
2. Model and Tokenizer:
 - We load the DialoGPT-medium model, which is a fine-tuned version of GPT-2 specifically
3. designed for conversations.
4. The tokenizer is responsible for encoding the input text into tokens that the model can understand and decoding the model's responses into human-readable text.

Chat Loop:

- The `chat_with_bot` function runs a loop where the user types a message and the bot generates a response.
- The loop continues until the user types "quit" to end the conversation.
- The user's input is encoded, and the model's response is generated using the `generate` method.
- The chat history (`chat_history_ids`) is maintained so that the model can generate contextually relevant responses (keeping the conversation coherent).

Text Generation:

- `model.generate` is used to produce text from the model. We set the `max_length` to limit the output size, and we use `temperature`, `no_repeat_ngram_size`, and other parameters to control the diversity of responses.
 - `temperature=0.7`: Controls randomness. Lower values (closer to 0) make the model
 - more deterministic.
- `no_repeat_ngram_size=3`: Prevents the model from repeating 3-grams (sequences of 3 words) in its output, which helps avoid repetitive responses.
- `top_p=0.9` and `top_k=50`: These parameters control the sampling strategy, which influences the diversity of generated responses.

Chat History:

- The chatbot remembers the entire conversation by appending the current user input to the chat history.
- This is important for maintaining context in conversations.

```
# Step 1: Import the necessary libraries
from transformers import GPT2LMHeadModel, GPT2Tokenizer
import torch

# Step 2: Load pre-trained DialoGPT model and tokenizer from Hugging Face
tokenizer = GPT2Tokenizer.from_pretrained("microsoft/DialoGPT-medium")
model = GPT2LMHeadModel.from_pretrained("microsoft/DialoGPT-medium")

# Step 3: Set the model in evaluation mode (important for inference)
model.eval()

# Step 4: Chat loop
def chat_with_bot():
    print("Chatbot: Hi, I'm here to chat with you! (Type 'quit' to end the conversation.)")

    # For storing chat history
    chat_history_ids = None
```

```

while True:
    # Take user input
    user_input = input("You: ")

    # Exit the chat if the user types 'quit'
    if user_input.lower() == 'quit':
        print("Chatbot: Goodbye!")
        break

    # Encode the user input and add the previous chat history (if any)
    new_user_input_ids = tokenizer.encode(user_input +
tokenizer.eos_token, return_tensors='pt')

    # Concatenate new user input with the previous chat history
    if chat_history_ids is not None:
        input_ids = torch.cat([chat_history_ids, new_user_input_ids],
dim=-1)
    else:
        input_ids = new_user_input_ids

    # Generate a response from the model
    chat_history_ids = model.generate(input_ids, max_length=1000,
pad_token_id=tokenizer.eos_token_id,
                                temperature=0.7,
no_repeat_ngram_size=3, top_p=0.9, top_k=50)

    # Decode and print the response (skip the user input part in the
output)
    bot_output = tokenizer.decode(chat_history_ids[:,
input_ids.shape[-1]:][0], skip_special_tokens=True)
    print(f"Chatbot: {bot_output}")

# Step 5: Start the chat
chat_with_bot()

```

OUTPUT

```
Chatbot: Hi, I'm here to chat with you! (Type 'quit' to end the
conversation.)
You: HI HOW ARE YOU
Chatbot: I'm good, how are you?
You: QUIT
Chatbot: Goodbye!
```

4. DEVELOPING A RECOMMENDATION SYSTEM USING COLLABORATIVE FILTERING OR DEEP LEARNING APPROACHES.

1. Collaborative Filtering

Collaborative filtering is based on the idea that users who agreed in the past will agree in the future. It predicts a user's preferences based on the preferences of other similar users.

There are two types of collaborative filtering:

- **User-based Collaborative Filtering:** Recommends items by finding similar users to the target user and recommending items those similar users liked.
- **Item-based Collaborative Filtering:** Recommends items similar to the items the user has already liked.

In this example, we'll focus on **Matrix Factorization**, which is a popular collaborative filtering technique.

Step 1: Install Libraries

```
pip install numpy pandas scikit-learn surprise
```

The `surprise` library is a Python library that implements collaborative filtering and matrix factorization methods.

Step 2: Prepare the Data

We'll use a **movie ratings dataset** as an example. You can use a dataset like **MovieLens**.

```
from surprise import Dataset, Reader
import pandas as pd

# Load the MovieLens dataset (example with ratings)
```

```

url =
"https://raw.githubusercontent.com/sidooms/MovieTweetings/master/latest/ratings
.dat"
ratings = pd.read_csv(url, sep="::", header=None, names=["user_id", "item_id",
"rating", "timestamp"])

# Prepare the data for surprise
reader = Reader(rating_scale=(1, 5))
data = Dataset.load_from_df(ratings[['user_id', 'item_id', 'rating']], reader)

```

- **ratings** contains user-item interactions, such as movie ratings.
- **Dataset.load_from_df()** prepares the data for use in the collaborative filtering model.

Step 3: Apply Collaborative Filtering using Matrix Factorization (SVD)

```

from surprise import SVD
from surprise.model_selection import train_test_split
from surprise import accuracy

# Split the dataset into train and test
trainset, testset = train_test_split(data, test_size=0.2)

# Use Singular Value Decomposition (SVD)
svd = SVD()

# Train the model
svd.fit(trainset)

# Test the model
predictions = svd.test(testset)

# Evaluate the accuracy
accuracy.rmse(predictions)

```

- **SVD (Singular Value Decomposition):** A matrix factorization technique used in collaborative filtering. It decomposes the user-item interaction matrix into latent factors, which represent hidden relationships between users and items.

Step 4: Make Recommendations

To make a recommendation for a specific user, we can predict ratings for all items and suggest the ones with the highest predicted ratings.

```

def recommend(user_id, n=10):
    # Get a list of all item IDs
    all_items = ratings['item_id'].unique()

    # Predict ratings for each item

```

```

predictions = [svd.predict(user_id, item_id) for item_id in all_items]

# Sort predictions by rating (descending order)
sorted_predictions = sorted(predictions, key=lambda x: x.est, reverse=True)

# Get the top n recommended items
top_n = sorted_predictions[:n]
return [(pred.iid, pred.est) for pred in top_n]

# Example: Recommend 10 items for user 1
recommendations = recommend(user_id=1, n=10)
print(recommendations)

```

OUTPUT:

```

<ipython-input-6-c85cc19332fc>:6: ParserWarning: Falling back to the
'python' engine because the 'c' engine does not support regex separators
(separators > 1 char and different from '\s+' are interpreted as regex);
you can avoid this warning by specifying engine='python'.

```

```

    ratings = pd.read_csv(url, sep="::", header=None, names=["user_id",
"item_id", "rating", "timestamp"])
RMSE: 2.9586
[(114508, 5), (499549, 5), (1305591, 5), (1428538, 5), (75314, 5),
(102926, 5), (114369, 5), (118715, 5), (120737, 5), (208092, 5)]

```

5. IMPLEMENTING COMPUTER VISION PROJECT SUCH AS OBJECT DETECTION OR IMAGE SEGMENTATION.**Explanation:**

1. **Upload Image:** The program uses **files.upload()** to upload an image directly to Colab.
2. **Face Detection:** The code detects faces in the uploaded image using the same Haar Cascade Classifier.
3. **Drawing Rectangles:** The faces are outlined with blue rectangles.
4. **Displaying the Image:** Since OpenCV uses BGR color format by default, we convert the image to RGB for proper display using **matplotlib**.

```

import cv2
from google.colab import files
import numpy as np
import matplotlib.pyplot as plt

```

```
# Upload an image file from your local system
uploaded = files.upload()

# Read the uploaded image
image_path = next(iter(uploaded)) # Get the file name of the uploaded image
img = cv2.imread(image_path)

# Convert image to grayscale
gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)

# Load pre-trained Haar Cascade Classifier for face detection
face_cascade = cv2.CascadeClassifier(cv2.data.harcascades + 'haarcascade_frontalface_default.xml')

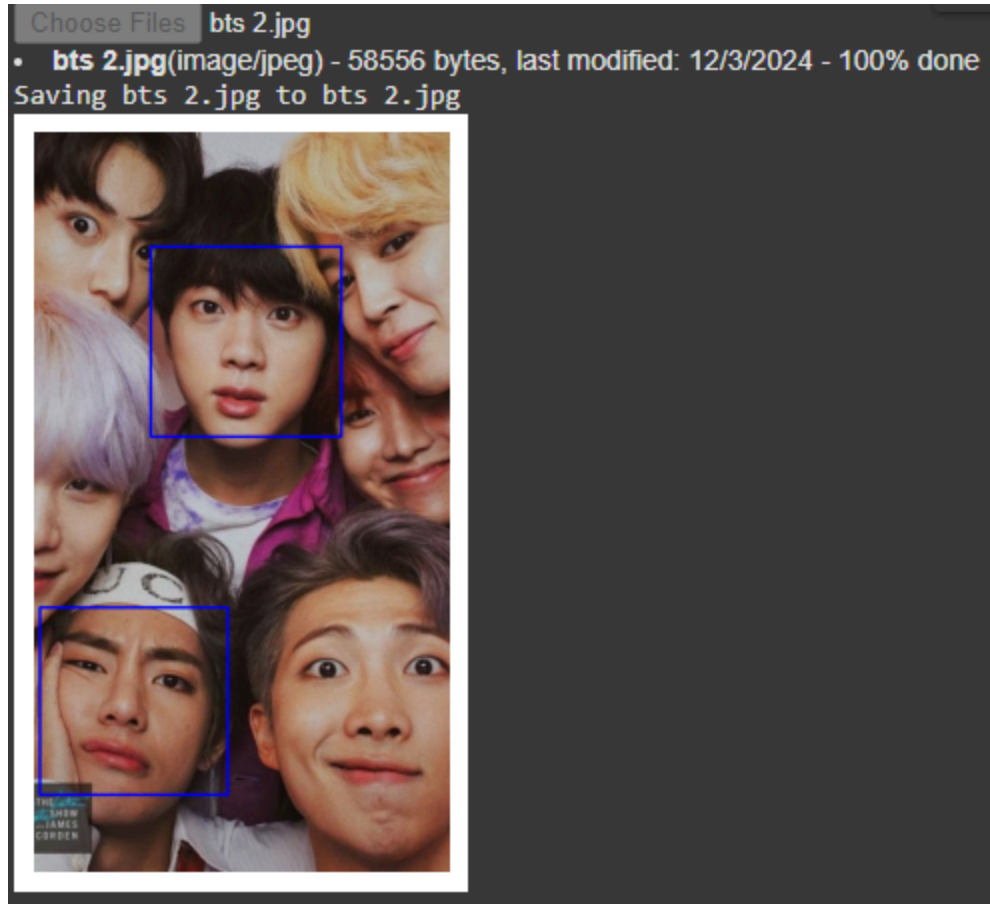
# Detect faces
faces = face_cascade.detectMultiScale(gray, scaleFactor=1.1, minNeighbors=5, minSize=(30, 30))

# Draw rectangles around the detected faces
for (x, y, w, h) in faces:
    cv2.rectangle(img, (x, y), (x + w, y + h), (255, 0, 0), 2)

# Convert the image from BGR to RGB for displaying in Colab
img_rgb = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)

# Display the image with detected faces
plt.imshow(img_rgb)
plt.axis('off') # Hide axis labels
plt.show()
```

OUTPUT



6. TRAINING A GENERATIVE ADVERSARIAL NETWORK (GAN) FOR GENERATING REALISTIC IMAGES

```
pip install tensorflow
```

```
import tensorflow as tf
from tensorflow.keras import layers
import numpy as np
import matplotlib.pyplot as plt
```

3. Load the Dataset

For this example, we will use the CIFAR-10 dataset (a dataset of 60,000 32x32 color images in 10 classes) for simplicity. You can substitute it with any other dataset.


```

# Load the CIFAR-10 dataset
(train_images, _), (_, _) = tf.keras.datasets.cifar10.load_data()

# Normalize images to range [-1, 1]
train_images = train_images.astype('float32')
train_images = (train_images - 127.5) / 127.5

# Reshape to the right shape (batch_size, height, width, channels)
train_images = train_images.reshape(train_images.shape[0], 32, 32, 3)

```

4. Build the Generator Model

The generator takes a random noise vector and generates an image.

```

def build_generator():
    model = tf.keras.Sequential()
    model.add(layers.Dense(256, input_dim=100)) # Input: random noise vector of size 100
    model.add(layers.LeakyReLU(alpha=0.2))
    model.add(layers.BatchNormalization(momentum=0.8))

    model.add(layers.Dense(512))
    model.add(layers.LeakyReLU(alpha=0.2))
    model.add(layers.BatchNormalization(momentum=0.8))

    model.add(layers.Dense(1024))
    model.add(layers.LeakyReLU(alpha=0.2))
    model.add(layers.BatchNormalization(momentum=0.8))

    model.add(layers.Dense(np.prod((32, 32, 3)), activation='tanh')) # Output image
    model.add(layers.Reshape((32, 32, 3)))

    return model

```

5. Build the Discriminator Model

The discriminator takes an image as input and outputs a probability (real or fake).

```

def build_discriminator():
    model = tf.keras.Sequential()

    model.add(layers.Flatten(input_shape=(32, 32, 3)))

```

```

model.add(layers.Dense(1024))

model.add(layers.LeakyReLU(alpha=0.2))


model.add(layers.Dense(512))

model.add(layers.LeakyReLU(alpha=0.2))


model.add(layers.Dense(256))

model.add(layers.LeakyReLU(alpha=0.2))


model.add(layers.Dense(1, activation='sigmoid')) # Output: real/fake

return model

```

6. Compile the Models

The generator is trained to minimize the discriminator's ability to classify the generated images correctly, while the discriminator is trained to maximize its ability to distinguish between real and fake images.

```

# Create the models
generator = build_generator()
discriminator = build_discriminator()

# Compile the discriminator
discriminator.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])

# The GAN model combines the generator and the discriminator
discriminator.trainable = False # Freeze the discriminator when training the GAN
gan_input = layers.Input(shape=(100,))
x = generator(gan_input)
gan_output = discriminator(x)
gan = tf.keras.Model(gan_input, gan_output)

```

```
# Compile the GAN
gan.compile(optimizer='adam', loss='binary_crossentropy')
```

7. Training the GAN

Now we define the training loop. For each batch of images:

- Train the discriminator on real and fake images.
- Train the generator via the GAN model to fool the discriminator

```
# Hyperparameters
```

```
epochs = 10000
```

```
batch_size = 64
```

```
half_batch = batch_size // 2
```

```
# Training loop
```

```
for epoch in range(epochs):
```

```
    # Train discriminator with real images
```

```
    idx = np.random.randint(0, train_images.shape[0], half_batch)
```

```
    real_images = train_images[idx]
```

```
    # Generate fake images
```

```
    noise = np.random.normal(0, 1, (half_batch, 100))
```

```
    fake_images = generator.predict(noise)
```

```
    # Train the discriminator (real images = 1, fake images = 0)
```

```
    d_loss_real = discriminator.train_on_batch(real_images, np.ones((half_batch, 1)))
```

```
d_loss_fake = discriminator.train_on_batch(fake_images, np.zeros((half_batch, 1)))

d_loss = 0.5 * np.add(d_loss_real, d_loss_fake)

# Train the generator (want to fool the discriminator)

noise = np.random.normal(0, 1, (batch_size, 100))

g_loss = gan.train_on_batch(noise, np.ones((batch_size, 1)))

# Print progress

if epoch % 1000 == 0:

    print(f'{epoch} [D loss: {d_loss[0]} | D accuracy: {100*d_loss[1]}] [G loss: {g_loss}]')

# Save generated images at intervals

if epoch % 1000 == 0:

    noise = np.random.normal(0, 1, (25, 100))

    generated_images = generator.predict(noise)

    generated_images = 0.5 * generated_images + 0.5 # Rescale to [0,1]

    fig, axs = plt.subplots(5, 5)

    count = 0

    for i in range(5):

        for j in range(5):

            axs[i, j].imshow(generated_images[count])

            axs[i, j].axis('off')

            count += 1
```

```
plt.show()
```

Key Points:

- **Training Loop:** The loop alternates between training the discriminator (on real and fake images) and training the generator (to fool the discriminator).
- **Image Generation:** Every 1000 epochs, the generated images are plotted to visualize progress.
- **Loss Functions:** Binary cross-entropy loss is used for both the generator and discriminator.

Output:

During training, you should see the generator loss decrease and the discriminator accuracy increase. Over time, the generator will produce more realistic images.

7. APPLYING REINFORCEMENT LEARNING ALGORITHMS TO SOLVE THE COMPLEX DECISION MAKING PROBLEMS.

Reinforcement Learning (RL) is a type of machine learning where an agent learns to make decisions by interacting with an environment to maximize cumulative rewards over time. RL is particularly well-suited for solving complex decision-making problems where the consequences of actions are delayed or uncertain.

Here's a high-level overview of how RL can be applied to solve such problems, along with an example using a popular RL algorithm called Q-learning.

Key Components of Reinforcement Learning:

1. **Agent:** The learner or decision maker.
2. **Environment:** The external system the agent interacts with.
3. **State:** A representation of the environment at a given time.
4. **Action:** A decision or move made by the agent.
5. **Reward:** Feedback from the environment in response to the agent's actions.
6. **Policy:** A strategy or mapping from states to actions that the agent uses to make decisions.
7. **Value Function:** A function that estimates how good a state is for the agent to be in.

Common RL Algorithms:

1. **Q-Learning:** A model-free algorithm that learns the value of action-state pairs (Q-values) and uses this to decide the best action in each state.

2. **Deep Q-Networks (DQN):** An extension of Q-learning using neural networks to approximate the Q-value function.
3. **Policy Gradient Methods:** Directly optimize the policy by adjusting its parameters based on feedback.
4. **Actor-Critic Methods:** Combine value-based and policy-based methods for more stable learning.

Example of Solving a Complex Decision-Making Problem using Q-Learning

Let's apply Q-learning to a gridworld environment, where an agent needs to navigate a grid to reach a goal while avoiding obstacles. This is a simple problem, but the principles extend to more complex decision-making tasks like robotics, games, and autonomous driving.

Steps:

1. **Define the Environment:** A 2D grid where the agent starts at one location and has to reach a goal, avoiding obstacles.
2. **Define the Reward Function:** The agent receives a positive reward when reaching the goal, and a negative reward for hitting obstacles or making unnecessary moves.
3. **Define the Q-table:** A table to store the Q-values, which represent the quality of each action in each state.
4. **Q-Learning Algorithm:** The agent will update the Q-values iteratively using the Bellman equation.

Step-by-Step Code Implementation:

Install Dependencies: You'll need **numpy** and **matplotlib** for this example. If you're running this locally, install them using:

```
pip install numpy matplotlib
```

1. **Code for Q-Learning in a Simple Gridworld:**

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
# Define the environment (4x4 grid)
```

```
grid_size = 4
```

```
goal_state = (3, 3)

start_state = (0, 0)

obstacles = [(1, 1), (2, 2)]


# Action space: Up, Down, Left, Right

actions = ['Up', 'Down', 'Left', 'Right']

action_map = {'Up': (-1, 0), 'Down': (1, 0), 'Left': (0, -1), 'Right': (0, 1)}


# Initialize the Q-table with zeros

Q = np.zeros((grid_size, grid_size, len(actions)))


# Define parameters

learning_rate = 0.1

discount_factor = 0.9

epsilon = 0.2 # Exploration factor

epochs = 1000


# Function to get next state based on action

def get_next_state(state, action):

    move = action_map[action]

    new_state = (state[0] + move[0], state[1] + move[1])


# Ensure the new state is within the grid boundaries
```

```
new_state = (max(0, min(grid_size-1, new_state[0])), max(0, min(grid_size-1, new_state[1])))
```

```
# If there's an obstacle, stay in the same position
```

```
if new_state in obstacles:
```

```
    return state
```

```
return new_state
```

```
# Function to choose an action (epsilon-greedy)
```

```
def choose_action(state):
```

```
    if np.random.rand() < epsilon:
```

```
        return np.random.choice(actions) # Exploration: Random action
```

```
    else:
```

```
        # Exploitation: Choose the action with the highest Q-value
```

```
        action_idx = np.argmax(Q[state[0], state[1]])
```

```
        return actions[action_idx]
```

```
# Function to run the Q-learning algorithm
```

```
def train_q_learning():
```

```
    for epoch in range(epochs):
```

```
        state = start_state
```

```
        total_reward = 0
```

```
        while state != goal_state:
```



```
    action = choose_action(state)

    next_state = get_next_state(state, action)


# Define the reward function

if next_state == goal_state:

    reward = 10 # Goal reached

elif next_state in obstacles:

    reward = -10 # Hit an obstacle

else:

    reward = -1 # Normal move


# Q-value update rule (Bellman Equation)

action_idx = actions.index(action)

Q[state[0], state[1], action_idx] += learning_rate * (reward + discount_factor *
np.max(Q[next_state[0], next_state[1]])) - Q[state[0], state[1], action_idx]


# Update state and accumulate reward

state = next_state

total_reward += reward


# Print progress every 100 episodes

if epoch % 100 == 0:

    print(f"Epoch {epoch}/{epochs}, Total Reward: {total_reward}")
```

```
# Train the agent using Q-learning
```

```
train_q_learning()
```

```
# Visualize the learned Q-values
```

```
fig, ax = plt.subplots(figsize=(8, 8))
```

```
for i in range(grid_size):
```

```
    for j in range(grid_size):
```

```
        # Skip the obstacles
```

```
        if (i, j) in obstacles:
```

```
            ax.text(j, i, "X", ha='center', va='center', color='red', fontsize=20)
```

```
        else:
```

```
            action_idx = np.argmax(Q[i, j]) # Choose action with max Q-value
```

```
            ax.text(j, i, actions[action_idx][0], ha='center', va='center', fontsize=20)
```

```
# Mark the start and goal positions
```

```
ax.text(start_state[1], start_state[0], "S", ha='center', va='center', color='green', fontsize=20)
```

```
ax.text(goal_state[1], goal_state[0], "G", ha='center', va='center', color='blue', fontsize=20)
```

```
ax.set_xticks(np.arange(grid_size))
```

```
ax.set_yticks(np.arange(grid_size))
```

```
ax.set_xticklabels([])
```

```
ax.set_yticklabels([])
```

```
ax.grid(True)
```

```
plt.show()
```

Explanation:

1. Environment Setup:

- A 4x4 grid is defined with a start state at (0, 0) and a goal state at (3, 3).
- There are obstacles at (1, 1) and (2, 2).
- The agent can move in four directions: Up, Down, Left, Right.

2. Q-Learning Algorithm:

- Q-Table: A table $Q[\text{state}, \text{action}]$ stores the Q-values, where each state-action pair has an associated value representing how good it is to take that action in that state.
- Exploration vs. Exploitation: The agent uses an epsilon-greedy policy to balance exploration (random actions) and exploitation (choosing the action with the highest Q-value).
- Q-Value Update: The Q-values are updated based on the Bellman equation, where the agent considers the reward it received and the future potential rewards.

3. Reward Structure:

- Positive reward (+10) for reaching the goal.
- Negative reward (-10) for hitting an obstacle.
- A small penalty (-1) for every other move to encourage efficiency.

4. Training Process:

- The agent takes actions in the environment, collects rewards, and updates its Q-values.
- After several epochs, the agent should have learned a policy that maximizes the cumulative reward by reaching the goal while avoiding obstacles.


5. Visualization:


- After training, the learned policy is displayed on the grid, showing the agent's best action for each state.

Output:

The output should show how the agent navigates the grid and learns the best actions to reach the goal while avoiding obstacles. The actions are represented by the first letter of each action (U, D, L, R) at each grid position.

```
Epoch 0/1000, Total Reward: -48  
Epoch 100/1000, Total Reward: 5  
Epoch 200/1000, Total Reward: 4  
Epoch 300/1000, Total Reward: 5  
Epoch 400/1000, Total Reward: 4  
Epoch 500/1000, Total Reward: 3  
Epoch 600/1000, Total Reward: 4  
Epoch 700/1000, Total Reward: 5  
Epoch 800/1000, Total Reward: 4  
Epoch 900/1000, Total Reward: 2
```





Applications of RL in Complex Decision-Making Problems:

Reinforcement Learning algorithms, like Q-learning, can be applied to more complex decision-making problems, such as:

1. **Robotics:** Teaching robots to navigate in an environment, pick and place objects, or perform tasks autonomously.
2. **Game AI:** Training agents to play games like chess, Go, or video games using reinforcement learning.
3. **Autonomous Vehicles:** Teaching self-driving cars to make decisions based on real-time traffic, road conditions, and safety.
4. **Finance:** Using RL for portfolio optimization, trading, or risk management in finance.
5. **Healthcare:** Training systems for personalized treatment recommendations, resource allocation in hospitals, or optimizing drug discovery processes.

In these complex problems, RL can help create agents capable of making high-level decisions by learning from the environment over time, balancing exploration and exploitation, and optimizing long-term objectives.

8. UTILISING TRANSFER LEARNING TO IMPROVE MODEL PERFORMANCE ON LIMITED DATASETS.

We will use Hugging Face for sentiment analysis. First, we will install transformers, if not already available in the environment:

```
!pip install transformers
```

In a typical NLP pipeline, several steps are taken to pre-process and prepare data to train and fine-tune the model. Hugging Face provides pipelines that perform all these steps with a few lines of code. Essentially, pipelines are composed of a tokeniser and a model to perform the task on the input text.

Below, we will utilise a pipeline for sentiment analysis using BERT on this text: “Transfer learning can help overcome issues related to over-reliance on data in machine learning”.

```
#importing pipeline
```

```
from transformers import pipeline
```

The pipeline will download and cache a default pre-trained model

(distilbert-base-uncased-finetuned-sst-2-english) and a tokeniser for sentiment analysis. We can then use the classifier on the input text:

```
# Initialising the pipeline with default model
```

```
classifier = pipeline('sentiment-analysis')
```

```
# Performing sentiment analysis on below text
```

```
classifier('Transfer learning can help overcome issues related
to over-reliance on data in machine learning')
```

The above pipeline produces a label “Positive” and a “score” of “0.908”. We can define

a model in the classifier and perform the analysis. Below we will use

“sentiment-roberta-large-english” [13].

```
#Initialising the pipeline with the selected model
```

```
classifier = pipeline('sentiment-analysis',
mode='siebert/sentiment-roberta-large-english')
```

```
# Performing sentiment analysis on below text
```

```
classifier('Transfer learning can help overcome issues related
to over-reliance on data in machine learning')
```

The above pipeline produces a label “Positive” and a “score” of “0.998”.

OUTPUT

```
[7] #Initialising the pipeline with the selected model
classifier = pipeline('sentiment-analysis', model='siebert/sentiment-roberta-large-english') # Cl
# Performing sentiment analysis on below text
classifier('Transfer learning can help overcome issues related to over-reliance on data in machir
```

File	Progress	Size	Speed
config.json	100%	687/687	9.44kB/s
pytorch_model.bin	100%	1.42G/1.42G	64.1MB/s
tokenizer_config.json	100%	256/256	7.80kB/s
vocab.json	100%	798k/798k	10.2MB/s
merges.txt	100%	456k/456k	13.3MB/s
special_tokens_map.json	100%	150/150	6.25kB/s

```
Device set to use cpu
[{'label': 'POSITIVE', 'score': 0.9986633062362671}]
```


9. BUILDING A DEEP LEARNING MODEL FOR TIME SERIES FORECASTING OR ANOMALY DETECTION.

Building a Deep Learning Model for Time Series Forecasting or Anomaly Detection

Time series forecasting and anomaly detection are crucial tasks in various fields like finance, healthcare, and industry. Deep learning models, such as **Recurrent Neural Networks (RNNs)**, **Long Short-Term Memory (LSTM)** networks, and **Gated Recurrent Units (GRUs)**, are particularly suited for time series tasks because they are designed to handle sequential data.

In this guide, we will:

1. **Build a deep learning model for time series forecasting** using **LSTM** (a type of RNN).
2. **Build a deep learning model for anomaly detection** in time series using an **autoencoder**.

Both models will use **TensorFlow** and **Keras**.

Part 1: Time Series Forecasting with LSTM

In time series forecasting, the goal is to predict future values of a time series based on past data.

Step 1: Install Dependencies

Install the necessary Python libraries if you haven't already:

```
pip install tensorflow numpy pandas matplotlib scikit-learn
```

Step 2: Load and Preprocess Data

We'll use the **Airline Passengers Dataset** as an example for time series forecasting. It contains monthly total passenger counts from 1949 to 1960. We will predict future passenger counts based on past data.

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.preprocessing import MinMaxScaler

# Load the dataset (you can replace this with any other time series dataset)
```

```

url =
"https://raw.githubusercontent.com/jbrownlee/Datasets/master/airline-passengers
.csv"
data = pd.read_csv(url, usecols=[1], engine='python', header=0)

# Visualize the data
plt.plot(data)
plt.title('Monthly Airline Passengers')
plt.xlabel('Month')
plt.ylabel('Passengers')
plt.show()

# Normalize the data (scaling to [0, 1] range for LSTM)
scaler = MinMaxScaler(feature_range=(0, 1))
data_scaled = scaler.fit_transform(data)

# Function to create the dataset in X (input) and y (output) for LSTM
def create_dataset(data, time_step=1):
    X, y = [], []
    for i in range(len(data)-time_step-1):
        X.append(data[i:(i+time_step), 0])
        y.append(data[i + time_step, 0])
    return np.array(X), np.array(y)

# Prepare the data for LSTM
time_step = 12 # Use 12 previous months to predict the next month
X, y = create_dataset(data_scaled, time_step)

# Reshape X for LSTM input: [samples, time steps, features]
X = X.reshape(X.shape[0], X.shape[1], 1)

```

Step 3: Build the LSTM Model

Now we define the **LSTM model** architecture:

```

from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM, Dense, Dropout

# Build the LSTM model
model = Sequential()
model.add(LSTM(units=50, return_sequences=False, input_shape=(X.shape[1], 1)))
model.add(Dropout(0.2)) # Dropout for regularization
model.add(Dense(units=1)) # Output layer (single value for forecasting)

# Compile the model
model.compile(optimizer='adam', loss='mean_squared_error')

```

Step 4: Train the Model

We will train the model on the time series data:

```
# Train the model
model.fit(X, y, epochs=20, batch_size=32)
```

Step 5: Make Predictions

Now that the model is trained, we can use it to make predictions:


```
# Predict the next 12 months
test_input = data_scaled[-time_step:] # Last 12 months of data for forecasting
test_input = test_input.reshape(1, -1, 1) # Reshape for LSTM input

# Predict
predicted = model.predict(test_input)
predicted = scaler.inverse_transform(predicted) # Inverse scaling

print(f'Predicted passengers for the next month: {predicted[0][0]}')
```

Step 6: Visualize the Predictions

```
# Plot the results
train_data = data[:len(data) - 12]
plt.plot(train_data, label='Training Data')
plt.plot(range(len(train_data), len(train_data) + 12), predicted,
label='Forecast', color='red')
plt.legend()
plt.title('Time Series Forecasting')
plt.xlabel('Month')
plt.ylabel('Passengers')
plt.show()
```



Part 2: Anomaly Detection with Autoencoders

Anomaly detection in time series is the task of identifying unusual patterns or outliers. An **autoencoder** is an unsupervised deep learning model used for anomaly detection. The autoencoder learns to compress (encode) the input into a latent representation and then reconstruct (decode) it. If the reconstruction error is high, the data point is likely an anomaly.

Step 1: Build the Autoencoder Model

We will build an autoencoder for anomaly detection using time series data.

```
from tensorflow.keras.layers import Input, Dense
from tensorflow.keras.models import Model

# Define the autoencoder model
input_dim = X.shape[1]
```

```

encoding_dim = 14 # Number of dimensions for the encoded representation

input_layer = Input(shape=(input_dim,))
encoded = Dense(encoding_dim, activation='relu')(input_layer)
decoded = Dense(input_dim, activation='sigmoid')(encoded)

# Create the autoencoder model
autoencoder = Model(inputs=input_layer, outputs=decoded)

# Compile the model
autoencoder.compile(optimizer='adam', loss='mean_squared_error')

# Train the autoencoder on the normal (non-anomalous) data
autoencoder.fit(X, X, epochs=50, batch_size=32)

```

Step 2: Detect Anomalies

We can now detect anomalies by checking the reconstruction error. If the error exceeds a certain threshold, the data point is considered anomalous.

```

# Make predictions using the autoencoder
reconstructed = autoencoder.predict(X)

# Calculate reconstruction error (Mean Squared Error)
reconstruction_error = np.mean(np.square(X.reshape(X.shape[0], X.shape[1])
- reconstructed), axis=1) # Reshape X again for subtraction

# Set a threshold for anomaly detection (this can be tuned)
threshold = np.percentile(reconstruction_error, 95)

# Identify anomalies (where reconstruction error is higher than threshold)
anomalies = reconstruction_error > threshold


# Visualize anomalies
plt.plot(data)

# Adjust the x-axis for the scatter plot to align with the original data
anomaly_indices = np.where(anomalies)[0] + time_step + 1 # Adjust for
time_step and the loop offset in create_dataset

plt.scatter(anomaly_indices, data.iloc[anomaly_indices, 0], color='red',
label='Anomalies') # Use iloc to access data by index

```

```
plt.title('Anomaly Detection in Time Series')
plt.xlabel('Time')
plt.ylabel('Value')
plt.legend()
plt.show()
```



Explanation of Key Steps

1. **LSTM for Time Series Forecasting:**
 - o **Data Preparation:** We prepared the time series data by scaling it and creating a sliding window of previous time steps (e.g., 12 months) to predict the next value.
 - o **Model Building:** The LSTM model learns from sequential patterns and is trained to predict future values.
 - o **Prediction:** After training, we used the LSTM model to forecast future values.
 2. **Autoencoder for Anomaly Detection:**
 - o **Autoencoder Architecture:** We defined an autoencoder network with an encoding and decoding layer. The model tries to reconstruct the input, and if the reconstruction error is high, we classify that as an anomaly.
 - o **Anomaly Detection:** Anomalies are detected by calculating the reconstruction error. If the error is above a certain threshold, it indicates an anomaly.
-

10. IMPLEMENTING A MACHINE LEARNING PIPELINE FOR AUTOMATED FEATURE ENGINEERING AND MODEL SELECTION.

A **Machine Learning Pipeline** automates various steps in the process of building and deploying machine learning models. The pipeline typically includes:

1. Data preprocessing and feature engineering
2. Model selection and evaluation
3. Hyperparameter tuning
4. Model training and deployment

In this guide, we'll build an **automated machine learning pipeline** that focuses on **automated feature engineering** and **model selection** using popular Python libraries like **scikit-learn**, **TPOT** (a tool for automated machine learning), and **Optuna** (a hyperparameter optimization library). The goal is to automate as many aspects of the machine learning workflow as possible, saving time and effort while ensuring optimal performance.

Key Components of the Pipeline

1. **Automated Feature Engineering:**
 - o This step focuses on automatically creating new features from existing data.
 - o It can include operations like scaling, encoding categorical variables, and handling missing values.
2. **Model Selection:**

- o Automating model selection involves testing multiple machine learning algorithms (e.g., decision trees, random forests, SVMs, etc.) to find the best-performing model for a given dataset.
3. **Hyperparameter Tuning:**
- o Tuning hyperparameters is crucial for optimizing model performance. This can be done automatically using methods like grid search or random search.

Step-by-Step Guide: Automated Machine Learning Pipeline

We'll demonstrate the pipeline using **scikit-learn** and the **TPOT** library for automated machine learning. **TPOT** helps automatically select the best features, models, and hyperparameters for your task.

Step 1: Install Dependencies

Install the necessary Python libraries:

```
pip install scikit-learn tpot optuna pandas numpy matplotlib
```

Step 2: Load and Preprocess Data

For the sake of illustration, we'll use the **Iris dataset** from `scikit-learn` (you can replace this with your own dataset):

```
import pandas as pd
from sklearn import datasets
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler

# Load the Iris dataset
iris = datasets.load_iris()
X = iris.data
y = iris.target

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
                                                    random_state=42)

# Feature scaling (important for algorithms like SVM and neural networks)
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

Step 3: Automated Feature Engineering (Optional)

While basic preprocessing like scaling is covered, more advanced feature engineering can be done using libraries like **Feature-engine** or **TPOT**. For this basic example, we'll focus on preprocessing using `scikit-learn`.

If you had more complex data (with categorical variables, missing values, etc.), automated feature engineering might involve:

- Encoding categorical variables (e.g., One-Hot Encoding, Label Encoding)
- Imputation of missing values
- Feature extraction techniques (e.g., PCA, polynomial features, etc.)

Step 4: Use TPOT for Automated Model Selection and Hyperparameter Tuning

TPOT uses genetic algorithms to automatically search for the best model and preprocessing pipeline. Here, we'll use TPOT to automate the selection of machine learning models and hyperparameters.

```
from tpot import TPOTClassifier

# Initialize the TPOTClassifier
tpot = TPOTClassifier( generations=5, population_size=20, random_state=42,
cv=5, verbosity=2)

# Train the model (TPOT automatically handles preprocessing and model
selection)
tpot.fit(X_train_scaled, y_train)

# Evaluate the model
accuracy = tpot.score(X_test_scaled, y_test)
print(f"Test accuracy: {accuracy:.4f}")

# Export the best pipeline
tpot.export('best_model_pipeline.py')
```

Step 5: Automated Hyperparameter Tuning with Optuna (Optional)

In addition to using TPOT, we can also use **Optuna** for hyperparameter optimization. Optuna allows you to define a search space and automatically search for the best hyperparameters using techniques like **Tree-structured Parzen Estimator (TPE)**.

Here is an example using **Optuna** to optimize hyperparameters for a **Random Forest Classifier**.

```
import optuna
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score
```

```

# Define the objective function for optimization
def objective(trial):
    # Hyperparameter space
    n_estimators = trial.suggest_int('n_estimators', 50, 200)
    max_depth = trial.suggest_int('max_depth', 5, 20)
    min_samples_split = trial.suggest_int('min_samples_split', 2, 10)

    # Train the model with the selected hyperparameters
    model = RandomForestClassifier(n_estimators=n_estimators,
max_depth=max_depth, min_samples_split=min_samples_split, random_state=42)
    model.fit(X_train_scaled, y_train)

    # Predict and calculate accuracy
    y_pred = model.predict(X_test_scaled)
    accuracy = accuracy_score(y_test, y_pred)
    return accuracy

# Create an Optuna study and optimize
study = optuna.create_study(direction="maximize")
study.optimize(objective, n_trials=20)

# Print the best hyperparameters found
print("Best hyperparameters:", study.best_params)

# Train the model with the best hyperparameters
best_params = study.best_params
best_model = RandomForestClassifier(**best_params, random_state=42)
best_model.fit(X_train_scaled, y_train)

# Evaluate the model
best_accuracy = best_model.score(X_test_scaled, y_test)
print(f"Best Model Test Accuracy: {best_accuracy:.4f}")

```

Step 6: Evaluate the Best Model

After training and tuning, you can evaluate the model's performance on the test set. Both **TPOT** and **Optuna** return models with optimized parameters, which you can then use to predict on new data.

```

# Evaluate the best model from TPOT or Optuna
y_pred = best_model.predict(X_test_scaled)
accuracy = accuracy_score(y_test, y_pred)
print(f"Test Accuracy of Final Model: {accuracy:.4f}")

```

Step 7: Model Deployment (Optional)

Once you have the best model, you can save it for future use or deployment using `joblib` or `pickle`.

```
import joblib

# Save the trained model
joblib.dump(best_model, 'best_model.pkl')

# Save the scaler for future use
joblib.dump(scaler, 'scaler.pkl')
```

This allows you to reload the model and use it for predictions on new, unseen data.

OUTPUT

```
[I 2025-01-12 03:44:33,519] A new study created in memory with name: no-name-a6d72daf-2e4a-4f62-af47-952f30dd4019
[I 2025-01-12 03:44:34,194] Trial 0 finished with value: 1.0 and parameters: {'n_estimators': 139, 'max_depth': 6, 'min_samples_split': 5}. Best is trial 0 with
[I 2025-01-12 03:44:34,616] Trial 1 finished with value: 1.0 and parameters: {'n_estimators': 133, 'max_depth': 14, 'min_samples_split': 3}. Best is trial 0 with
[I 2025-01-12 03:44:34,920] Trial 2 finished with value: 1.0 and parameters: {'n_estimators': 100, 'max_depth': 13, 'min_samples_split': 7}. Best is trial 0 with
[I 2025-01-12 03:44:35,428] Trial 3 finished with value: 1.0 and parameters: {'n_estimators': 137, 'max_depth': 6, 'min_samples_split': 3}. Best is trial 0 with
[I 2025-01-12 03:44:35,600] Trial 4 finished with value: 1.0 and parameters: {'n_estimators': 54, 'max_depth': 10, 'min_samples_split': 5}. Best is trial 0 with
[I 2025-01-12 03:44:36,317] Trial 5 finished with value: 1.0 and parameters: {'n_estimators': 171, 'max_depth': 15, 'min_samples_split': 7}. Best is trial 0 with
[I 2025-01-12 03:44:37,428] Trial 6 finished with value: 1.0 and parameters: {'n_estimators': 171, 'max_depth': 12, 'min_samples_split': 8}. Best is trial 0 with
[I 2025-01-12 03:44:37,869] Trial 7 finished with value: 1.0 and parameters: {'n_estimators': 75, 'max_depth': 18, 'min_samples_split': 9}. Best is trial 0 with
[I 2025-01-12 03:44:38,325] Trial 8 finished with value: 1.0 and parameters: {'n_estimators': 131, 'max_depth': 10, 'min_samples_split': 2}. Best is trial 0 with
[I 2025-01-12 03:44:38,579] Trial 9 finished with value: 1.0 and parameters: {'n_estimators': 57, 'max_depth': 17, 'min_samples_split': 7}. Best is trial 0 with
[I 2025-01-12 03:44:38,910] Trial 10 finished with value: 1.0 and parameters: {'n_estimators': 200, 'max_depth': 5, 'min_samples_split': 5}. Best is trial 0 with
[I 2025-01-12 03:44:39,104] Trial 11 finished with value: 1.0 and parameters: {'n_estimators': 108, 'max_depth': 8, 'min_samples_split': 4}. Best is trial 0 with
[I 2025-01-12 03:44:39,388] Trial 12 finished with value: 1.0 and parameters: {'n_estimators': 155, 'max_depth': 15, 'min_samples_split': 2}. Best is trial 0 with
[I 2025-01-12 03:44:39,591] Trial 13 finished with value: 1.0 and parameters: {'n_estimators': 104, 'max_depth': 20, 'min_samples_split': 4}. Best is trial 0 with
[I 2025-01-12 03:44:39,845] Trial 14 finished with value: 1.0 and parameters: {'n_estimators': 146, 'max_depth': 8, 'min_samples_split': 5}. Best is trial 0 with
[I 2025-01-12 03:44:40,061] Trial 15 finished with value: 1.0 and parameters: {'n_estimators': 123, 'max_depth': 13, 'min_samples_split': 3}. Best is trial 0 with
[I 2025-01-12 03:44:40,366] Trial 16 finished with value: 1.0 and parameters: {'n_estimators': 164, 'max_depth': 11, 'min_samples_split': 10}. Best is trial 0 with
[I 2025-01-12 03:44:40,714] Trial 17 finished with value: 1.0 and parameters: {'n_estimators': 199, 'max_depth': 15, 'min_samples_split': 6}. Best is trial 0 with
[I 2025-01-12 03:44:40,926] Trial 18 finished with value: 1.0 and parameters: {'n_estimators': 117, 'max_depth': 7, 'min_samples_split': 3}. Best is trial 0 with
[I 2025-01-12 03:44:41,091] Trial 19 finished with value: 1.0 and parameters: {'n_estimators': 91, 'max_depth': 9, 'min_samples_split': 4}. Best is trial 0 with
Best hyperparameters: {'n_estimators': 139, 'max_depth': 6, 'min_samples_split': 5}
Best Model Test Accuracy: 1.0000
```

```
# Evaluate the best model from TPOT or Optuna
y_pred = best_model.predict(X_test_scaled)
accuracy = accuracy_score(y_test, y_pred)
print(f"Test Accuracy of Final Model: {accuracy:.4f}")

Test Accuracy of Final Model: 1.0000

7] import joblib

# Save the trained model
joblib.dump(best_model, 'best_model.pkl')

# Save the scaler for future use
joblib.dump(scaler, 'scaler.pkl')

['scaler.pkl']
```

11. USING ADVANCED OPTIMIZATION TECHNIQUES LIKE EVOLUTIONARY ALGORITHMS OR BAYESIAN OPTIMIZATION FOR HYPERPARAMETER TUNING.

Certainly! Let's go over how to perform hyperparameter tuning using Evolutionary Algorithms and Bayesian Optimization in Google Colab. These are advanced optimization techniques that can significantly improve your model's performance.

We'll implement:

1. Evolutionary Algorithms using the TPOT library.
2. Bayesian Optimization using the Optuna library.

1. Evolutionary Algorithms with TPOT

TPOT (Tree-based Pipeline Optimization Tool) is an AutoML library that uses genetic algorithms to optimize machine learning pipelines, which includes hyperparameter tuning. It's one of the evolutionary algorithms where the model is optimized over generations.

Install TPOT (if not already installed)

```
!pip install tpot
```

Code Example using TPOT for Hyperparameter Tuning

We'll use the Iris dataset for classification, but this approach can be adapted to any dataset.

```
# Import necessary libraries
```

```
import pandas as pd
```

```
import numpy as np
```

```
from sklearn.datasets import load_iris
```

```
from sklearn.model_selection import train_test_split
```

```
from tpot import TPOTClassifier

from sklearn.metrics import accuracy_score


# Load the Iris dataset

data = load_iris()

df = pd.DataFrame(data.data, columns=data.feature_names)

df['target'] = data.target


# Split the data into features (X) and target (y)

X = df.drop('target', axis=1)

y = df['target']


# Split the data into training and testing sets

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)


# Initialize the TPOTClassifier for model optimization with
evolutionary algorithms

tpot = TPOTClassifier( generations=5, population_size=20,
random_state=42, verbosity=2)


# Fit the model on the training data

tpot.fit(X_train, y_train)
```

```
# Evaluate the model on the test set

y_pred = tpot.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)

print(f"Accuracy of the optimized model: {accuracy}")

# Export the best pipeline found by TPOT

tpot.export('best_model_pipeline.py')
```

OUTPUT

```
is_regressor
is_classifier
is_regressor
is_classifier
is_classifier
is_classifier
is_regressor
is_classifier
is_regressor
is_classifier

Generation 1 - Current best internal CV score: 0.975
Generation 2 - Current best internal CV score: 0.975
Generation 3 - Current best internal CV score: 0.975
Generation 4 - Current best internal CV score: 0.975
Generation 5 - Current best internal CV score: 0.975

Best pipeline: MLPClassifier(input_matrix, alpha=0.0001, learning_rate_init=0.001)
Accuracy of the optimized model: 1.0
```

Explanation:

1. `TPOTClassifier`: Optimizes the pipeline using genetic algorithms.
 - `generations`: Number of generations to run the optimization (higher means more search, but takes longer).
 - `population_size`: Number of models evaluated per generation.
2. `Evaluation`: After training, we evaluate the model on the test set and print the accuracy.
3. `Export`: TPOT exports the best model pipeline to a Python file for later use.

2. Bayesian Optimization with Optuna

Optuna is a hyperparameter optimization framework that performs Bayesian optimization to find the optimal set of hyperparameters for machine learning models. It efficiently explores the hyperparameter space by making intelligent decisions based on previous evaluations.

Install Optuna (if not already installed)

```
!pip install optuna
```

Code Example using Optuna for Hyperparameter Tuning

We'll demonstrate Bayesian Optimization for Random Forest using Optuna.

```
import optuna

import pandas as pd

import numpy as np

from sklearn.datasets import load_iris

from sklearn.model_selection import train_test_split

from sklearn.ensemble import RandomForestClassifier
```

```
from sklearn.metrics import accuracy_score

# Load the Iris dataset

data = load_iris()

df = pd.DataFrame(data.data, columns=data.feature_names)

df['target'] = data.target


# Split the data into features (X) and target (y)

X = df.drop('target', axis=1)

y = df['target']


# Split the data into training and testing sets

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)


# Objective function for Optuna

def objective(trial):

    # Define hyperparameters to optimize

    n_estimators = trial.suggest_int('n_estimators', 10, 200)

    max_depth = trial.suggest_int('max_depth', 2, 20)

    min_samples_split = trial.suggest_int('min_samples_split',
2, 10)
```



```
min_samples_leaf = trial.suggest_int('min_samples_leaf', 1,
10)

# Initialize the model with suggested hyperparameters
model = RandomForestClassifier(n_estimators=n_estimators,
                              max_depth=max_depth,

min_samples_split=min_samples_split,

min_samples_leaf=min_samples_leaf,

                              random_state=42)

# Train the model
model.fit(X_train, y_train)

# Make predictions and evaluate accuracy
y_pred = model.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)

return accuracy
```

```
# Create the study object for Bayesian Optimization

study = optuna.create_study(direction='maximize')


# Optimize hyperparameters

study.optimize(objective, n_trials=50)


# Print the best hyperparameters and accuracy

print(f"Best hyperparameters: {study.best_params}")

print(f"Best accuracy: {study.best_value}")
```

Explanation:

1. Objective Function: Defines the function to optimize, where `trial.suggest_int()` is used to define hyperparameters (such as `n_estimators`, `max_depth`, etc.). Optuna will explore different combinations of hyperparameters.
2. Study: We create a study with the direction `'maximize'`, which means we aim to maximize the accuracy.
3. Optimization: We perform the optimization using `study.optimize()`, and it will run `n_trials` number of times, testing different combinations of hyperparameters.
4. Best Parameters: After the optimization process, we print the best hyperparameters and corresponding accuracy.

Key Benefits of Bayesian Optimization with Optuna:

- Efficiency: Bayesian optimization efficiently explores the hyperparameter space by learning from previous evaluations.
- Scalability: It works well for both small and large parameter spaces.

Conclusion

Both Evolutionary Algorithms (using TPOT) and Bayesian Optimization (using Optuna) provide powerful methods for optimizing machine learning models. Here's a summary of when to use each:

1. TPOT:

- Best for AutoML tasks, where you want the system to automatically select and tune models.
- Uses genetic algorithms to optimize the whole pipeline, including feature engineering and model selection.

2. Optuna:

- Best for fine-tuning specific hyperparameters of a model using Bayesian optimization.
- Efficient for searching hyperparameter spaces, especially in scenarios where you have predefined models.

Both approaches are highly effective in automating hyperparameter tuning and improving model performance, making them ideal for tasks where optimal model configuration is crucial.

12. DEPLOYING A MACHINE LEARNING MODEL IN A PRODUCTION ENVIRONMENT USING CONTAINERIZATION AND CLOUD SERVICES.

Deploying a machine learning model into production involves several steps, including model training, packaging the model, containerization, and deployment using cloud services. Here, we'll walk you through a simple yet effective pipeline to deploy a machine learning model using containerization (Docker) and cloud services (e.g., Google Cloud, AWS, or Azure).

For the sake of simplicity, let's break it into smaller steps :

1. Train a Machine Learning Model (using the Iris dataset).
2. Save the Trained Model to a file (e.g., using `joblib`).
3. Containerize the Model using Docker (with a simple API interface).
4. Deploy to Cloud (we will illustrate using Google Cloud, but the steps are similar for other cloud platforms).

Step 1: Train a Model in Google Colab

We'll start by training a simple Random Forest Classifier model on the Iris dataset.

```
# Import necessary libraries
```

```
import pandas as pd
```

```
from sklearn.datasets import load_iris
```

```
from sklearn.ensemble import RandomForestClassifier
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.metrics import accuracy_score
```

```
import joblib # For saving the model
```

```
# Load Iris dataset
```

```
data = load_iris()
```

```
df = pd.DataFrame(data.data, columns=data.feature_names)
```

```
df['target'] = data.target
```

```
# Split the data into features (X) and target (y)
```

```
X = df.drop('target', axis=1)

y = df['target']


# Split the data into training and testing sets

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)


# Initialize and train the model

model = RandomForestClassifier(random_state=42)

model.fit(X_train, y_train)


# Make predictions and evaluate the model

y_pred = model.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)

print(f"Model accuracy: {accuracy}")


# Save the model using joblib

joblib.dump(model, 'iris_model.pkl')
```

Step 2: Prepare the Model for Deployment

Once the model is trained, we save it using **joblib**, which is a simple way to serialize the model. The saved model (**iris_model.pkl**) will be used for deployment.

Step 3: Dockerize the Model

To deploy the model in a production environment, we will create a Docker container. Docker allows us to package the model and its dependencies into a single container, which can then be run in any environment.

Creating a Flask API for the Model

We will build a simple Flask web service that will serve the trained model.

1. Create a Flask Application: We will create a Python script (`app.py`) that will load the model and handle predictions via a web API.

```
# Create a simple Flask app (this would typically be in app.py)
```

```
from flask import Flask, request, jsonify
```

```
import joblib
```

```
import numpy as np
```

```
# Initialize Flask app
```

```
app = Flask(__name__)
```

```
# Load the trained model (this assumes the model is in the  
current directory)
```

```
model = joblib.load('iris_model.pkl')
```

```
# Define a route for predictions
```

```
@app.route('/predict', methods=['POST'])

def predict():

    try:

        # Get the data from the POST request

        data = request.get_json(force=True)

        # Convert the input data into an array

        input_data = np.array(data['features']).reshape(1, -1)

        # Make prediction

        prediction = model.predict(input_data)

        # Return the result

        return jsonify({'prediction': prediction[0]})

    except Exception as e:

        return jsonify({'error': str(e)})

# Run the app

if __name__ == '__main__':

    app.run(debug=True)
```

Creating a Dockerfile

Next, we create a Dockerfile to containerize the model and Flask app. Here's a simple **Dockerfile** that sets up the container.

```
# Use an official Python runtime as a parent image
```

```
FROM python:3.8-slim
```

```
# Set the working directory in the container
```

```
WORKDIR /app
```

```
# Copy the current directory contents into the container at /app
```

```
COPY . /app
```

```
# Install any needed packages specified in requirements.txt
```

```
RUN pip install --no-cache-dir -r requirements.txt
```

```
# Expose port 5000 for the Flask app
```

```
EXPOSE 5000
```

```
# Run app.py when the container launches
```

```
CMD ["python", "app.py"]
```

Creating requirements.txt

This file lists the necessary Python packages to be installed in the container.

flask

joblib

scikit-learn

numpy

Step 4: Build and Run the Docker Container Locally

1. Build the Docker Image: Run the following command to build the Docker image locally.

```
docker build -t iris-model-app .
```

2. Run the Docker Container: After building the image, run the container.

```
docker run -p 5000:5000 iris-model-app
```

Now your model is deployed locally and accessible at <http://localhost:5000>. You can test it by sending POST requests to the `/predict` endpoint.

Step 5: Deploy to Google Cloud (or Any Cloud Service)

Once the Docker image is built, you can deploy it to a cloud service like Google Cloud, AWS, or Azure. We'll demonstrate deployment using Google Cloud Run, which allows you to run Docker containers in a fully managed environment.

Deploy to Google Cloud Run:

1. Install Google Cloud SDK: You need the Google Cloud SDK to interact with Google Cloud from your terminal. Install it if you haven't already: [Google Cloud SDK Installation](#).

Authenticate with Google Cloud: Authenticate your account using the command:

```
gcloud auth login
```

2.

Create a Google Cloud Project (if you don't have one):

```
gcloud projects create YOUR_PROJECT_ID
```

3.

Push Docker Image to Google Container Registry:

First, tag your Docker image:

```
docker tag iris-model-app gcr.io/YOUR_PROJECT_ID/iris-model-app
```

Then, push the image to Google Container Registry:

```
docker push gcr.io/YOUR_PROJECT_ID/iris-model-app
```

4.

Deploy the Image on Google Cloud Run:

Once the image is pushed to the registry, you can deploy it to Google Cloud Run using:

```
gcloud run deploy --image gcr.io/YOUR_PROJECT_ID/iris-model-app  
--platform managed
```

5. This command will deploy your container and provide a URL for your deployed application. You can now access your

Flask API via the URL provided.

Conclusion

In this tutorial, we have covered:

1. Training a machine learning model (Random Forest on Iris dataset).
2. Saving the model to a file using `joblib`.
3. Containerizing the model using Docker, and creating a Flask API to serve the model.
4. Deploying the model to Google Cloud Run (or any other cloud service).

With this pipeline, you can deploy machine learning models into production environments efficiently using containerization and cloud services. You can adapt the same steps to other models, containers, and cloud platforms as well.

13. USE PYTHON LIBRARIES SUCH AS GPT-2 OR TEXTGENRM TO TRAIN GENERATIVE MODELS ON A CORPUS OF TEXT DATA AND GENERATE NEW TEXT BASED ON PATTERNS IT HAS LEARNED.

Training generative models such as GPT-2 or TextGenRnn in Python can be a powerful approach to generate new text based on patterns learned from a given corpus of text. Below, we'll walk through how to use the GPT-2 model for text generation in Google Colab.

We'll cover:

1. Using GPT-2 with Transformers for generating text.
2. Training a generative model (TextGenRnn) using a custom dataset.

1. Text Generation with GPT-2 in Google Colab

For this section, we will use the pre-trained GPT-2 model from the Transformers library by Hugging Face.

Step 1: Install Required Libraries

In Colab, start by installing the required libraries, including Transformers and Torch.

```
!pip install transformers
```

```
!pip install torch
```

Step 2: Load Pre-trained GPT-2 Model

We'll use the pre-trained GPT-2 model for text generation. GPT-2 is already trained on a vast corpus of text, so you can use it for text generation out of the box.

```
import torch
```

```
from transformers import GPT2LMHeadModel, GPT2Tokenizer
```

```
# Load pre-trained GPT-2 model and tokenizer
```

```
model_name = 'gpt2' # You can use 'gpt2-medium' or 'gpt2-large' for larger models
```

```
model = GPT2LMHeadModel.from_pretrained(model_name)
```

```
tokenizer = GPT2Tokenizer.from_pretrained(model_name)
```

```
# Make the model ready for inference
```

```
model.eval()
```

```
# Device configuration (Use GPU if available)
```

```
device = torch.device("cuda" if torch.cuda.is_available() else
"cpu")

model.to(device)


# Generate text function

def generate_text(prompt, max_length=100):

    inputs = tokenizer.encode(prompt,
return_tensors='pt').to(device)


    # Generate text using GPT-2 model

    outputs = model.generate(inputs, max_length=max_length,
num_return_sequences=1, no_repeat_ngram_size=2, top_p=0.95,
temperature=0.7, do_sample=True)


    # Decode and return the generated text

    text = tokenizer.decode(outputs[0],
skip_special_tokens=True)

    return text


# Example usage

prompt = "Once upon a time in a land far away"

generated_text = generate_text(prompt, max_length=100)

print(generated_text)
```

Explanation:

- **GPT2LMHeadModel**: The GPT-2 model that we are using for language modeling.
- **GPT2Tokenizer**: The tokenizer used to encode and decode text.
- **generate_text**: A function that generates text based on a prompt. You can control the **max_length**, **top_p**, and **temperature** for varied results.
- **Device**: We ensure the model runs on the GPU if available for faster generation.

2. Training a Generative Text Model Using TextGenRnn

Now, let's look at training a generative text model with TextGenRnn. TextGenRnn is a simple-to-use library for training a recurrent neural network (RNN) on your custom text data.

Step 1: Install TextGenRnn

First, install the textgenrnn library. This is an easy-to-use Python library for text generation.

```
!pip install textgenrnn
```

Step 2: Prepare Your Custom Dataset

Before training a model, you need a dataset of text that the model will learn from. You can upload your own text corpus or use an example corpus. Let's assume you have a dataset stored in a file (e.g., **your_text_data.txt**).

```
# If you want to upload your own dataset
```

```
from google.colab import files
```

```
uploaded = files.upload()
```

```
# Check the file
```

```
!cat your_text_data.txt
```

Alternatively, for testing purposes, we can use a simple example text corpus:

```
# Simple example dataset for text generation
```

```
text_data = """
```

```
Once upon a time, in a village far away, there lived a young  
girl named Ella.
```

```
Ella loved to explore the forest and discover new things. Every  
day, she would
```

```
wander through the woods, discovering hidden secrets and meeting  
new friends.
```

```
One day, she stumbled upon an old tree that was said to grant  
wishes...
```

```
"""
```

```
# Save the data to a text file (you can skip this if you already  
have a dataset)
```

```
with open('your_text_data.txt', 'w') as file:
```

```
    file.write(text_data)
```

Step 3: Train the TextGenRnn Model

Now, let's train the model using the TextGenRnn library.

```
from textgenrnn import textgenrnn

# Initialize the text generation model
textgen = textgenrnn.TextgenRnn()

# Train the model on your dataset
textgen.train_from_file('your_text_data.txt', num_epochs=10)

# Save the trained model
textgen.save()

# Generate text using the trained model
generated_text = textgen.generate(return_as_list=True)[0]
print(generated_text)
```

Explanation:

- `textgenrnn.TextgenRnn()`: Initializes the text generation model.
- `train_from_file()`: Trains the model using your custom dataset (`your_text_data.txt`).
- `num_epochs=10`: You can adjust the number of epochs based on your data size and desired training time.
- `generate()`: Generates text using the trained model.

3. Summary of How It Works

- GPT-2: Uses a pre-trained transformer model that has been trained on massive corpora and can generate text based on patterns it has learned. You can use it directly in Colab with minimal setup.
- TextGenRnn: Trains an RNN-based model on your custom text data and generates text based on that learned pattern.

14. EXPERIMENT WITH NEURAL NETWORKS LIKE GAN'S USING PYTHON LIBRARIES LIKE TENSORFLOW OR PYTORCH TO GENERATE NEW IMAGES BASED ON A DATASET OF IMAGES.

Generating new images using Generative Adversarial Networks (GANs) is a popular approach in machine learning. We'll use TensorFlow and Keras in Google Colab to experiment with GANs and generate new images based on a dataset of images. For simplicity, we'll use the Fashion MNIST dataset, which contains images of clothing, to train a simple GAN.

Steps:

1. Install TensorFlow (if it's not already installed).
2. Load the Fashion MNIST dataset.
3. Build and train a simple GAN.
4. Generate new images based on the learned distribution.

Step 1: Install TensorFlow

If you don't have TensorFlow installed in your Colab environment, you can install it using:

```
!pip install tensorflow
```

Step 2: Load the Fashion MNIST Dataset

We will use the Fashion MNIST dataset, which contains 60,000 training images and 10,000 test images of 28x28 grayscale images of clothing items.

```
import tensorflow as tf

from tensorflow.keras.datasets import fashion_mnist

import matplotlib.pyplot as plt

import numpy as np

# Load Fashion MNIST dataset

(x_train, _), (_, _) = fashion_mnist.load_data()

# Normalize the images to [0, 1]

x_train = x_train / 255.0

# Reshape images to (batch_size, 28, 28, 1) to match the input
of a CNN

x_train = np.expand_dims(x_train, axis=-1)

# Display a sample image

plt.imshow(x_train[0].reshape(28, 28), cmap='gray')

plt.show()
```

Step 3: Build the GAN

A GAN consists of two neural networks:

- Generator: Generates fake images from random noise.
- Discriminator: Tries to distinguish between real and fake images.

We'll use Keras to build both the generator and discriminator models.

Generator Model

The generator takes random noise as input and generates an image.

```
from tensorflow.keras.models import Sequential

from tensorflow.keras.layers import Dense, Reshape, Flatten,
BatchNormalization, LeakyReLU
```

```
def build_generator(latent_dim):

    model = Sequential()

    model.add(Dense(128, input_dim=latent_dim))

    model.add(LeakyReLU(0.2))

    model.add(BatchNormalization(momentum=0.8))

    model.add(Dense(256))

    model.add(LeakyReLU(0.2))

    model.add(BatchNormalization(momentum=0.8))

    model.add(Dense(512))

    model.add(LeakyReLU(0.2))
```

```

model.add(BatchNormalization(momentum=0.8))

model.add(Dense(1024))

model.add(LeakyReLU(0.2))

model.add(BatchNormalization(momentum=0.8))

model.add(Dense(np.prod((28, 28, 1)), activation='tanh'))

model.add(Reshape((28, 28, 1)))

return model

```

Discriminator Model

The discriminator is a binary classifier that predicts whether an image is real or fake.

```

from tensorflow.keras.layers import Conv2D, Conv2DTranspose

```

```

def build_discriminator():

    model = Sequential()

    model.add(Conv2D(64, kernel_size=3, strides=2,
padding='same', input_shape=(28, 28, 1)))

    model.add(LeakyReLU(0.2))

    model.add(Conv2D(128, kernel_size=3, strides=2,
padding='same'))

    model.add(LeakyReLU(0.2))

    model.add(Conv2D(256, kernel_size=3, strides=2,
padding='same'))

```

```
model.add(LeakyReLU(0.2))

model.add(Flatten())

model.add(Dense(1, activation='sigmoid'))

return model
```

Combined Model (GAN)

We will also create the combined model, where the generator is trained to fool the discriminator.

```
from tensorflow.keras.models import Model
```

```
def build_gan(generator, discriminator):

    discriminator.trainable = False

    model = Sequential()

    model.add(generator)

    model.add(discriminator)

    return model
```

Step 4: Compile the Models

Now we compile the models. We'll use binary crossentropy for both the discriminator and GAN models, and Adam optimizer for optimization.

```
from tensorflow.keras.optimizers import Adam
```

```
latent_dim = 100 # Dimensionality of the random noise input to
the generator
```

```
optimizer = Adam(learning_rate=0.0002, beta_1=0.5)
```

```
# Build and compile the discriminator
```

```
discriminator = build_discriminator()
```

```
discriminator.compile(loss='binary_crossentropy',
optimizer=optimizer, metrics=['accuracy'])
```

```
# Build the generator
```

```
generator = build_generator(latent_dim)
```

```
# Build the GAN model
```

```
gan = build_gan(generator, discriminator)
```

```
gan.compile(loss='binary_crossentropy', optimizer=optimizer)
```

Step 5: Training the GAN

Now we define the training loop. We will train the generator to produce better fake images that the discriminator cannot distinguish from real images.

```
def train_gan(generator, discriminator, gan, x_train, epochs,
batch_size, latent_dim):
```

```
    # Labels for real and fake images
```

```
    real_labels = np.ones((batch_size, 1))
```

```
fake_labels = np.zeros((batch_size, 1))

for epoch in range(epochs):

    # Train the discriminator

    idx = np.random.randint(0, x_train.shape[0], batch_size)

    real_images = x_train[idx]

    noise = np.random.normal(0, 1, (batch_size, latent_dim))

    fake_images = generator.predict(noise)

    # Train on real images and fake images

    d_loss_real = discriminator.train_on_batch(real_images,
real_labels)

    d_loss_fake = discriminator.train_on_batch(fake_images,
fake_labels)

    d_loss = 0.5 * np.add(d_loss_real, d_loss_fake)

    # Train the generator via the GAN model (we want to fool
the discriminator)

    noise = np.random.normal(0, 1, (batch_size, latent_dim))

    g_loss = gan.train_on_batch(noise, real_labels)
```

```
if epoch % 1000 == 0:

    print(f"Epoch: {epoch}, D Loss: {d_loss[0]}, G Loss:
{g_loss}")

    # Save generated images to visualize progress

    if epoch % 5000 == 0:

        plot_generated_images(epoch, generator,
latent_dim)

def plot_generated_images(epoch, generator, latent_dim,
examples=10, dim=(1, 10), figsize=(10, 1)):

    noise = np.random.normal(0, 1, (examples, latent_dim))

    generated_images = generator.predict(noise)

    plt.figure(figsize=figsize)

    for i in range(examples):

        plt.subplot(dim[0], dim[1], i+1)

        plt.imshow(generated_images[i].reshape(28, 28),
cmap='gray')

        plt.axis('off')

    plt.tight_layout()

    plt.savefig(f"generated_image_{epoch}.png")

    plt.close()
```



```
# Train the GAN
```

```
train_gan(generator, discriminator, gan, x_train, epochs=10000,  
batch_size=64, latent_dim=latent_dim)
```

Step 6: Generate New Images

After training, we can generate new images based on random noise. The generator will have learned the distribution of the Fashion MNIST images and will generate new images resembling the ones it was trained on.

```
# Generate new images using the trained generator
```

```
def generate_images(generator, latent_dim, num_images=10):
```

```
    noise = np.random.normal(0, 1, (num_images, latent_dim))
```

```
    generated_images = generator.predict(noise)
```

```
    for i in range(num_images):
```

```
        plt.imshow(generated_images[i].reshape(28, 28),  
cmap='gray')
```

```
        plt.axis('off')
```

```
        plt.show()
```

```
# Generate new images
```

```
generate_images(generator, latent_dim)
```

Conclusion

We:

1. Built a Generative Adversarial Network (GAN) with TensorFlow/Keras.
2. Trained the GAN using the Fashion MNIST dataset.
3. Generated new images based on the patterns the GAN learned.

The GAN learns to generate images that resemble those in the training dataset, and as training progresses, the quality of generated images improves.

```

# Practical 1A - Load a CSV dataset. Handle missing
values, inconsistent formatting,
# and outliers.
import pandas as pd
import numpy as np
# Create a dataset
data = {
    'Name': ['Alice', 'Bob', 'Charlie', 'David', 'Eve',
'Frank', 'Grace', 'Heidi', np.nan, 'Ivan'],
    'Age': [25, 32, np.nan, 45, 28, 35, 29, None, 40, 38],
    'Salary': [50000, 60000, 70000, np.nan, 50000, 120000,
45000, 52000, 680000, 40000], # Contains outliers
    'Date of Joining':
['2022-01-10', '2021-05-15', '2020/02/20', '2019-10-05', '202
2-03-30', '2021-07-10',

'2018-11-01', '2019-12-31', '2020-04-15', '2021-08-01']
}
# Convert to a DataFrame
df = pd.DataFrame(data)
print(df)
# Save to a CSV file
df.to_csv('dataset1a.csv', index=False)
df.head()
# Display basic information about the dataset
print(df.info()) # Shows column names, non-null counts,
and data types
print(df.head()) # Displays the first few rows

# Standardize date formatting
df['Date of Joining'] = pd.to_datetime(df['Date of
Joining'], errors='coerce') # Converts to datetime
format
# Check for any remaining NaT values in the date column
(if conversion fails)

```

```

df.dropna(subset=['Date of Joining'], inplace=True)  #
Remove rows with invalid dates

# Calculate Q1 and Q3 for Salary
Q1 = df['Salary'].quantile(0.25)
Q3 = df['Salary'].quantile(0.75)
IQR = Q3 - Q1

# Define the upper and lower bounds for detecting
outliers
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR
# Filter out outliers in the 'Salary' column
df = df[(df['Salary'] >= lower_bound) & (df['Salary'] <=
upper_bound)]

print(df.info())  # Check for any remaining missing
values
print(df.describe())  # Get summary statistics
print(df)  # Display the cleaned dataset

df['Age'] = df['Age'].fillna(0)
print(df)  # Display the cleaned dataset

```

```

Name, Age, Salary, Date of Joining
Alice, 25.0, 50000.0, 2022-01-10
Bob, 32.0, 60000.0, 2021-05-15
Charlie, , 70000.0, 2020/02/20
David, 45.0, , 2019-10-05
Eve, 28.0, 50000.0, 2022-03-30
Frank, 35.0, 120000.0, 2021-07-10
Grace, 29.0, 45000.0, 2018-11-01
Heidi, , 52000.0, 2019-12-31
, 40.0, 680000.0, 2020-04-15
Ivan, 38.0, 40000.0, 2021-08-01

```

Output 👍

```

      Name  Age  Salary Date of Joining
0  Alice  25.0  50000.0   2022-01-10
1   Bob   32.0  60000.0   2021-05-15
2 Charlie  NaN  70000.0   2020/02/20
3  David  45.0    NaN   2019-10-05
4   Eve   28.0  50000.0   2022-03-30
5  Frank  35.0 120000.0   2021-07-10
6  Grace  29.0  45000.0   2018-11-01
7  Heidi  NaN  52000.0   2019-12-31
8   NaN   40.0 680000.0   2020-04-15
9   Ivan  38.0  40000.0   2021-08-01

```

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 10 entries, 0 to 9

Data columns (total 4 columns):

```
#   Column          Non-Null Count  Dtype
---

```

```
0  Name          9 non-null    object
1  Age           8 non-null    float64
2  Salary        9 non-null    float64
3  Date of Joining 10 non-null   object
```

dtypes: float64(2), object(2)

memory usage: 452.0+ bytes

None

```

      Name  Age  Salary Date of Joining
0  Alice  25.0  50000.0   2022-01-10
1   Bob   32.0  60000.0   2021-05-15
2 Charlie  NaN  70000.0   2020/02/20
3  David  45.0    NaN   2019-10-05
4   Eve   28.0  50000.0   2022-03-30

```

<class 'pandas.core.frame.DataFrame'>

Index: 6 entries, 0 to 9

Data columns (total 4 columns):

```
#   Column          Non-Null Count  Dtype
---

```

```
0  Name          6 non-null    object
1  Age           5 non-null    float64
2  Salary        6 non-null    float64
3  Date of Joining 6 non-null    datetime64[ns]
```

dtypes: datetime64[ns](1), float64(2), object(1)

memory usage: 240.0+ bytes

None

	Age	Salary	Date of Joining
count	5.000000	6.000000	6
mean	30.400000	49500.000000	2021-01-03 08:00:00
min	25.000000	40000.000000	2018-11-01 00:00:00
25%	28.000000	46250.000000	2020-05-04 06:00:00
50%	29.000000	50000.000000	2021-06-23 00:00:00
75%	32.000000	51500.000000	2021-11-30 12:00:00
max	38.000000	60000.000000	2022-03-30 00:00:00
std	4.929503	6745.368782	NaN

	Name	Age	Salary	Date of Joining
0	Alice	25.0	50000.0	2022-01-10
1	Bob	32.0	60000.0	2021-05-15
4	Eve	28.0	50000.0	2022-03-30
6	Grace	29.0	45000.0	2018-11-01
7	Heidi	NaN	52000.0	2019-12-31
9	Ivan	38.0	40000.0	2021-08-01

	Name	Age	Salary	Date of Joining
0	Alice	25.0	50000.0	2022-01-10
1	Bob	32.0	60000.0	2021-05-15
4	Eve	28.0	50000.0	2022-03-30
6	Grace	29.0	45000.0	2018-11-01
7	Heidi	0.0	52000.0	2019-12-31
9	Ivan	38.0	40000.0	2021-08-01

Process finished with exit code 0

```
#Practical 1B - Load a dataset, calculate descriptive  
summary statistics, create visualizations using different  
graphs,  
# and identify potential features and target variables  
#Note: Explore Univariate and Bivariate graphs  
(Matplotlib)  
# and Seaborn for visualization.
```

```
import matplotlib.pyplot as plt  
import pandas as pd  
import seaborn as sns  
from sklearn import datasets
```

```
x = [1,2,3,4,5,6,7,8]  
y = [0.1,0.2,0.5,0.6,0.7,0.8,0.9,1.3]  
plt.plot(x,y)  
plt.show()
```

```
# initialise data of lists  
data = { 'Name':[ 'AAA', 'BBB', 'CCC', 'DDD'],  
         'Age':[ 30,21,29,28],  
         'Weight':[60,65,50,48]}  
df = pd.DataFrame(data)  
print(df.describe())
```

```
#Box Plot  
sns.boxplot(df,x=data['Age'],y=data['Weight'])  
plt.show()
```

```
#Line Plot  
sns.lineplot(df,x=data['Age'],y=data['Weight'])  
plt.show()
```

```
#Violin Plot  
sns.violinplot(data['Age'])  
plt.show()
```

```

#Iris Dataset Plotting - KDE Plot
# Setting up the Data Frame
iris = datasets.load_iris()
iris_df = pd.DataFrame(iris.data,
columns=['Sepal_Length',
          'Sepal_Width',
          'Patal_Length', 'Petal_Width'])
iris_df['Target'] = iris.target
iris_df['Target'].replace([0], 'Iris_Setosa',
inplace=True)
iris_df['Target'].replace([1], 'Iris_Vercicolor',
inplace=True)
iris_df['Target'].replace([2], 'Iris_Virginica',
inplace=True)
# Plotting the KDE Plot
sns.kdeplot(iris_df.loc[(iris_df['Target'] ==
'Iris_Virginica'),
'Sepal_Length'], color='b', fill=True)

plt.xlabel('Data')
plt.ylabel('Density')
plt.title('Univariate KDE Plot')
# Show the plot
plt.show()

#Practical 1C : Create or Explore datasets to use all
#pre-processing routines like label encoding, scaling,
and binarization.
import pandas as pd
from sklearn.preprocessing import LabelEncoder
#create DataFrame
df = pd.DataFrame({'team': ['A', 'A', 'B', 'B', 'B', 'C',
'C', 'D'],

```




```


        'position': ['G', 'F', 'G', 'F', 'F',
'G', 'G', 'F'],
        'all_star': ['Y', 'N', 'Y', 'Y', 'Y',
'N', 'Y', 'N'],
        'points': [11, 8, 10, 6, 6, 5, 9, 12]})
#view DataFrame
print(df)
#perform label encoding across team, position, and
all_star columns
df[['team', 'position', 'all_star']] = df[['team', 'position',
'all_star']].apply(LabelEncoder().fit_transform)
#view updated DataFrame
print(df)

```


Output 👍


	Age	Weight
count	4.000000	4.000000
mean	27.000000	55.750000
std	4.082483	8.098354
min	21.000000	48.000000
25%	26.250000	49.500000
50%	28.500000	55.000000
75%	29.250000	61.250000
max	30.000000	65.000000






Machine Learning







```

#Practical 1C : Create or Explore datasets to use all
#pre-processing routines like label encoding, scaling,
and binarization.
import pandas as pd
from sklearn.preprocessing import LabelEncoder
#create DataFrame
df = pd.DataFrame({'team': ['A', 'A', 'B', 'B', 'B', 'C',
                             'C', 'D'],
                   'position': ['G', 'F', 'G', 'F', 'F',
                                'G', 'G', 'F'],
                   'all_star': ['Y', 'N', 'Y', 'Y', 'Y',
                                'N', 'Y', 'N'],
                   'points': [11, 8, 10, 6, 6, 5, 9, 12]})
#view DataFrame
print(df)
#perform label encoding across team, position, and
all_star columns
df[['team', 'position', 'all_star']] = df[['team', 'position',
                                             'all_star']].apply(LabelEncoder().fit_transform)
#view updated DataFrame
print(df)

```

Output 👍

	team	position	all_star	points
0	A	G	Y	11
1	A	F	N	8
2	B	G	Y	10
3	B	F	Y	6
4	B	F	Y	6
5	C	G	N	5
6	C	G	Y	9
7	D	F	N	12

	team	position	all_star	points
0	0	1	1	11
1	0	0	0	8
2	1	1	1	10

3	1	0	1	6
4	1	0	1	6
5	2	1	0	5
6	2	1	1	9
7	3	0	0	12

Process finished with exit code 0

```
#Practical 2 - Implement and demonstrate the FIND-S
algorithm
# for finding the most specific hypothesis based on a
given set
# of training data samples. Read the training data from a
CSV file
# and generate the final specific hypothesis. (Create
your dataset)
```

```
import csv
num_attributes=7
a=[]

with open('PractData2.csv', 'r') as csvfile:
    reader = csv.reader(csvfile)
    next(reader)
    for row in reader:
        a.append(row)
        print(row)

print("\n The initial value of hypothesis: ")
hypothesis = ['0'] * num_attributes
print(hypothesis)

for j in range(0,num_attributes):
    hypothesis[j] = a[0][j];
```



```

print("\n Find S: Finding a Maximally Specific
Hypothesis\n")
for i in range(0,len(a)):
    if a[i][num_attributes]=='Yes':
        for j in range(0,num_attributes):
            if a[i][j]!=hypothesis[j]:
                hypothesis[j]='?'
            else :
                hypothesis[j]= a[i][j]
    print(" For Training instance No:{0} the hypothesis is
".format(i),hypothesis)
print("\n The Maximally Specific Hypothesis for a given
Training Examples :\n")
print(hypothesis)

```

```

Example ,Sky ,AirTemp ,Humidity ,Wind, Water ,Forecast
,EnjoySport
1,Sunny,Warm,Normal,Strong,Warm,Same,Yes
2,Sunny,Warm,High,Strong,Warm,Same,Yes
3,Rainy,Cold,High,Strong,Warm,Change,No
4,Sunny,Warm,High,Strong,Cool,Change,Yes

```

Output 👍

```

['1', 'Sunny', 'Warm', 'Normal', 'Strong', 'Warm', 'Same', 'Yes']
['2', 'Sunny', 'Warm', 'High', 'Strong', 'Warm', 'Same', 'Yes']
['3', 'Rainy', 'Cold', 'High', 'Strong', 'Warm', 'Change', 'No']
['4', 'Sunny', 'Warm', 'High', 'Strong', 'Cool', 'Change', 'Yes']

```

The initial value of hypothesis:
['0', '0', '0', '0', '0', '0', '0']

Find S: Finding a Maximally Specific Hypothesis

For Training instance No:0 the hypothesis is ['1', 'Sunny', 'Warm', 'Normal', 'Strong', 'Warm', 'Same']

For Training instance No:1 the hypothesis is ['?', 'Sunny', 'Warm', '?', 'Strong', 'Warm', 'Same']

For Training instance No:2 the hypothesis is ['?', 'Sunny', 'Warm', '?', 'Strong', 'Warm', 'Same']
 For Training instance No:3 the hypothesis is ['?', 'Sunny', 'Warm', '?', 'Strong', '?', '?']

The Maximally Specific Hypothesis for a given Training Examples :

['?', 'Sunny', 'Warm', '?', 'Strong', '?', '?']

Process finished with exit code 0

```
#Practical 3A - Simple Linear Regression
#Fit a linear regression model on a dataset. Interpret
coefficients,
#make predictions, and evaluate performance using metrics
like R-squared and MSE
from statistics import linear_regression

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.model_selection import train_test_split

#Creating Testing Data
np.random.seed(0)
X=np.random.rand(100,1)
Y= 20*X+1+np.random.randn(100,1)

#Convert The above Data into DataFrame
data=pd.DataFrame({'X':X.flatten(), 'Y':Y.flatten()})
plt.scatter(X,Y)
plt.xlabel('Independent Variable (X)')
plt.ylabel('Dependent variable (Y)')
plt.title('Sample Data For Regression')
plt.show()
```

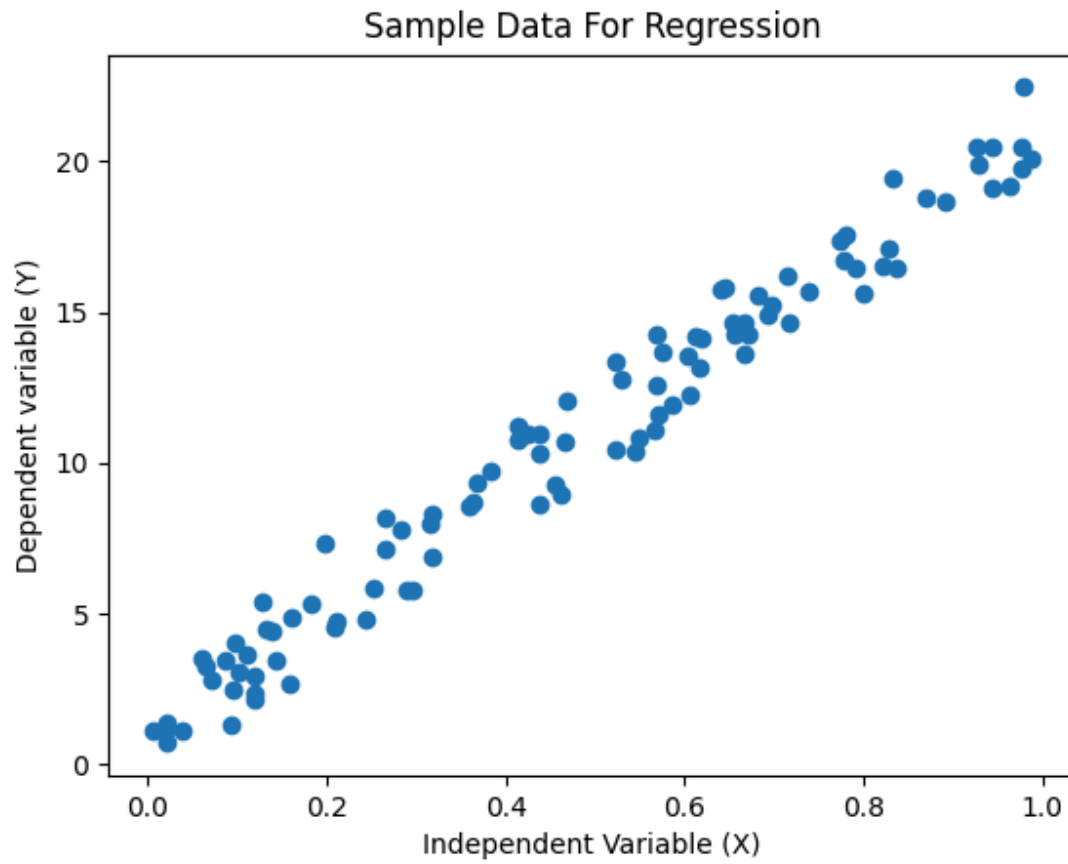
```
#Creating Training and Testing Data
X_train,X_test,Y_train,Y_test=train_test_split(X,Y,test_s
ize=0.2,random_state=42)
model=LinearRegression()
model.fit(X_train,Y_train)
y_pred=model.predict(X_test)
#Calculate MSE
mse=mean_squared_error(Y_test,y_pred)
#Calculate R-Squared
r_squared=r2_score(Y_test,y_pred)

print(f'Mean Squared Error (MSE):{mse:.2f}')
print(f'R-squared (rsq):{r_squared:.2f}')
```

Output 👍

Mean Squared Error (MSE):0.92
R-squared (rsq):0.96

Process finished with exit code 0



#Practical 3B - Extend linear regression to multiple features. Handle feature selection and #potential multicollinearity.

```
import numpy as np
import matplotlib as mpl
import matplotlib.pyplot as plt
from mpl_toolkits.mplot3d import Axes3D

def generate_dataset(n):
    x = []
    y = []
    random_x1 = np.random.rand()
    random_x2 = np.random.rand()
    for i in range(n):
        x1 = i
        x2 = i/2 + np.random.rand()*n
        x.append([1, x1, x2])
        y.append(random_x1 * x1 + random_x2 * x2 + 1)
    return np.array(x), np.array(y)

x, y = generate_dataset(200)


mpl.rcParams['legend.fontsize'] = 12

fig = plt.figure()
ax = fig.add_subplot(projection='3d')

ax.scatter(x[:, 1], x[:, 2], y, label='y', s=5)
ax.legend()
ax.view_init(45, 0)

plt.show()
```

Output 👍



```
#Practical 3C - Implement regression variants like  
# LASSO and Ridge on any generated dataset
```

```
import pandas as pd  
from sklearn.model_selection import train_test_split  
from sklearn.linear_model import LinearRegression  
from sklearn.linear_model import Lasso  
from sklearn.metrics import r2_score  
  
df=pd.read_csv('PractData3C.csv')  
df.head()
```

```

df.info()
df.drop(['model'], axis = 1, inplace = True)
print(df)
#create feature and Target Variable
x = df.drop(columns = 'hp', axis = 1)
y = df['hp']

#splitting the data into train and test set
x_train,x_test,y_train,y_test =
train_test_split(x,y,test_size=0.3,random_state=4)

#fitting data into the model
lr = LinearRegression()
print(lr.fit(x_train, y_train))

#r2-score for train data
x_pred_train = lr.predict(x_train)
print(r2_score(y_train, x_pred_train))

lasso = Lasso()
print(lasso.fit(x_train, y_train))

x_pred_lasso_test = lasso.predict(x_test)
print(r2_score(y_test, x_pred_lasso_test))

```

```

model,mpg,cyl,disp,hp
Honda,21,6,160,110
Toyota,22,6,180,90
Fiat,23,4,240,100
Audi,21,6,220,110

```

Output 👍

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4 entries, 0 to 3

```

Data columns (total 5 columns):

Column Non-Null Count Dtype

```

---
0  model  4 non-null  object
1  mpg    4 non-null  int64
2  cyl    4 non-null  int64
3  disp   4 non-null  int64
4  hp     4 non-null  int64

```

dtypes: int64(4), object(1)

memory usage: 292.0+ bytes

```

   mpg  cyl  disp  hp
0   21   6   160  110
1   22   6   180   90
2   23   4   240  100
3   21   6   220  110

```

LinearRegression()

1.0

Lasso()

-11.093699999999995

#Practical 4A - Perform binary classification using logistic regression.

Calculate accuracy, precision, recall, and understand the ROC curve.

```
import matplotlib.pyplot as plt
```

```
# roc curve and auc
```

```
from sklearn.datasets import make_classification
```

```
from sklearn.linear_model import LogisticRegression
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.metrics import roc_curve
```

```
from sklearn.metrics import roc_auc_score
```

```
# generate 2 class dataset
```

```
X, y = make_classification(n_samples=1000, n_classes=2,
random_state=1)
```

```
# split into train/test sets
```



```

trainX, testX, trainy, testy = train_test_split(X, y,
test_size=0.5, random_state=2)
# generate a no skill prediction (majority class)
ns_probs = [0 for _ in range(len(testy))]
# fit a model
model = LogisticRegression(solver='lbfgs')
model.fit(trainX, trainy)
# predict probabilities
lr_probs = model.predict_proba(testX)
# keep probabilities for the positive outcome only
lr_probs = lr_probs[:, 1]

# calculate scores
ns_auc = roc_auc_score(testy, ns_probs)
lr_auc = roc_auc_score(testy, lr_probs)

# summarize scores
print('No Skill: ROC AUC=%.3f' % (ns_auc))
print('Logistic: ROC AUC=%.3f' % (lr_auc))


# calculate roc curves
ns_fpr, ns_tpr, _ = roc_curve(testy, ns_probs)
lr_fpr, lr_tpr, _ = roc_curve(testy, lr_probs)
# plot the roc curve for the model
plt.plot(ns_fpr, ns_tpr, linestyle='--', label='No
Skill')
plt.plot(lr_fpr, lr_tpr, marker='.', label='Logistic')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.legend()
plt.show()

```

Output 👍

No Skill: ROC AUC=0.500
Logistic: ROC AUC=0.903

Process finished with exit code 0



```

#Practical 4B - Implement and demonstrate k-nearest
Neighbor algorithm.
# Read the training data
#from a .CSV file and build the model to classify a test
sample. Print both correct
#and wrong predictions.

# Import necessary modules
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import train_test_split
from sklearn.datasets import load_iris
import numpy as np
import matplotlib.pyplot as plt

# Loading data
irisData = load_iris()
# Create feature and target arrays
X = irisData.data
y = irisData.target
# Split into training and test set
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size = 0.2, random_state=42)
knn = KNeighborsClassifier(n_neighbors=7)
knn.fit(X_train, y_train)
# Predict on dataset which model has not seen before
print(knn.predict(X_test))
# Calculate the accuracy of the model
print(knn.score(X_test, y_test))
irisData = load_iris() # Create feature and target arrays
X = irisData.data
y = irisData.target
# Split into training and test set
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42)
neighbors = np.arange(1, 9)
train_accuracy = np.empty(len(neighbors))

```

```

test_accuracy = np.empty(len(neighbors))

# Loop over K values
for i, k in enumerate(neighbors):
    knn = KNeighborsClassifier(n_neighbors=k)
    knn.fit(X_train, y_train)
    # Compute training and test data accuracy
    train_accuracy[i] = knn.score(X_train, y_train)
    test_accuracy[i] = knn.score(X_test, y_test)

# Generate plot
plt.plot(neighbors, test_accuracy, label='Testing dataset
Accuracy')
plt.plot(neighbors, train_accuracy, label='Training
dataset Accuracy')

plt.legend()
plt.xlabel('n_neighbors')
plt.ylabel('Accuracy')
plt.show()

```


Output 👍

```

[1 0 2 1 1 0 1 2 2 1 2 0 0 0 0 1 2 1 1 2 0 2 0 2 2 2 2 0 0]
0.9666666666666667

```

Process finished with exit code 0



```
#Practical 4C - Build a decision tree classifier or  
regressor. Control hyperparameters like tree  
#depth to avoid overfitting. Visualize the tree  
from sklearn.datasets import load_iris  
from sklearn.tree import DecisionTreeClassifier  
from sklearn.model_selection import train_test_split  
from sklearn.metrics import accuracy_score  
  
# Load the iris dataset  
iris = load_iris()  
X = iris.data  
y = iris.target  
  
# Split the data into training and testing sets  
X_train, X_test, y_train, y_test = train_test_split(X, y,  
test_size=0.3)  
  
# Train a decision tree on the training set  
clf = DecisionTreeClassifier()  
clf.fit(X_train, y_train)  
  
# Predict on the training set  
y_train_pred = clf.predict(X_train)  
train_accuracy = accuracy_score(y_train, y_train_pred)  
  
# Predict on the testing set  
y_test_pred = clf.predict(X_test)  
test_accuracy = accuracy_score(y_test, y_test_pred)  
  
# Compare the accuracy on the training set and the  
testing set  
print("Accuracy on the training set:", train_accuracy)  
print("Accuracy on the testing set:", test_accuracy)  
  
from sklearn.tree import DecisionTreeClassifier
```

```
from sklearn.datasets import make_classification

X, y = make_classification(n_samples=1000, n_features=20,
                           n_classes=2)
clf = DecisionTreeClassifier(random_state=0)
clf.fit(X, y)

path = clf.cost_complexity_pruning_path(X, y)
alphas, impurities = path.ccp_alphas, path.impurities
print(alphas)
```

Output 👍

Accuracy on the training set: 1.0

Accuracy on the testing set: 0.9777777777777777

```
[0.      0.00096703 0.00098113 0.00098667 0.00098958 0.00099448
 0.00132317 0.00171429 0.00176471 0.0018221  0.00184615 0.00196078
 0.00269779 0.0027033  0.00285714 0.0032967  0.00365008 0.00405809
 0.00667198 0.02607097 0.32000808]
```

```
#Practical 4D -Implement a Support Vector Machine for any  
relevant dataset
```

```
# Load the important packages
```

```
from sklearn.datasets import load_iris  
import matplotlib.pyplot as plt  
from sklearn.inspection import DecisionBoundaryDisplay  
from sklearn.svm import SVC
```

```
# Load the datasets
```

```
iris = load_iris()  
X = iris.data[:, :2]  
y = iris.target
```

```
#Build the model
```

```
svm = SVC(kernel="rbf", gamma=0.5, C=1.0)
```

```
# Trained the model
```

```
svm.fit(X, y)
```

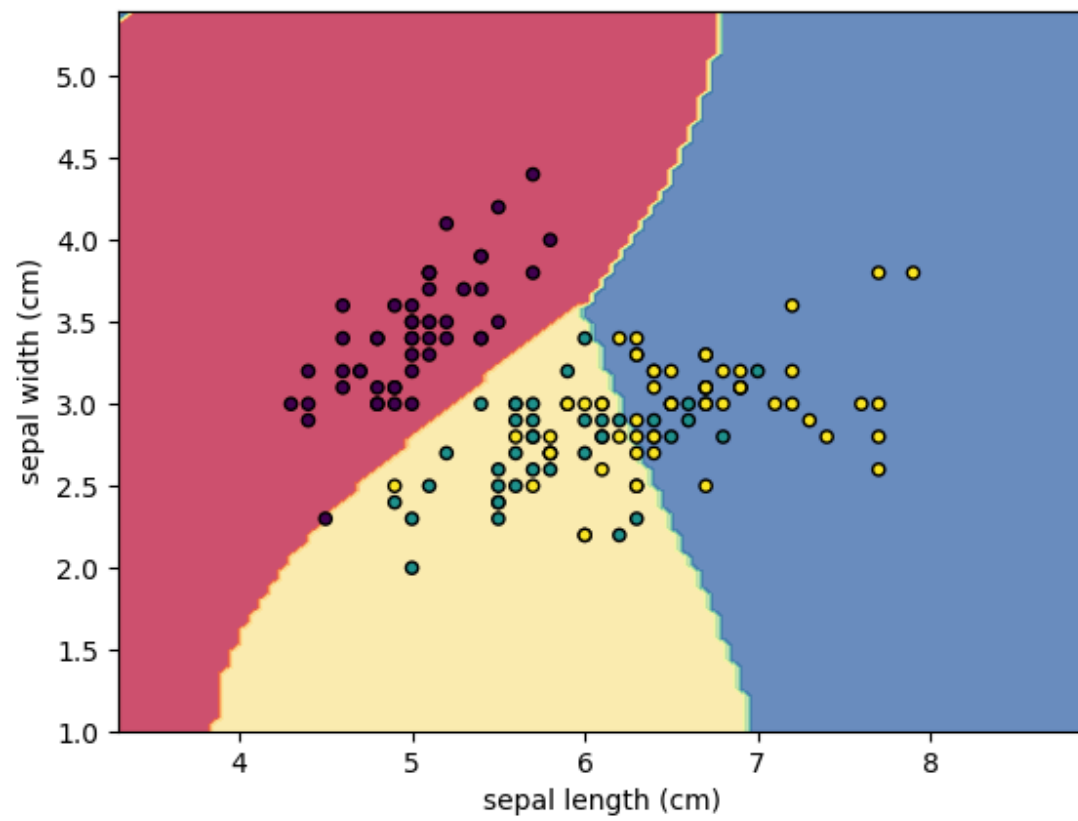
```
# Plot Decision Boundary
```

```
DecisionBoundaryDisplay.from_estimator(  
    svm,  
    X,  
    response_method="predict",  
    cmap=plt.cm.Spectral,  
    alpha=0.8,  
    xlabel=iris.feature_names[0],  
    ylabel=iris.feature_names[1],  
    )
```

```
# Scatter plot
```

```
plt.scatter(X[:, 0], X[:, 1],  
            c=y,  
            s=20, edgecolors="k")  
plt.show()
```


Output 👍



```
#Practical 4E- Train a random forest ensemble.
#Experiment with the number of trees and
#feature sampling. Compare performance to a single
decision tree.

# Import necessary libraries
import pandas as pd
from sklearn.datasets import fetch_california_housing
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error, r2_score

# Load the California Housing dataset
california_housing = fetch_california_housing()
california_data = pd.DataFrame(california_housing.data,

columns=california_housing.feature_names)
#print(california_data)
california_data['MEDV'] = california_housing.target

# Select relevant features and target variable
X = california_data.drop('MEDV', axis=1)
y = california_data['MEDV']

# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2,

random_state=42)
# Create a Random Forest Regressor
rf_regressor = RandomForestRegressor(n_estimators=100,
random_state=42)
# Train the regressor
rf_regressor.fit(X_train, y_train)
# Make predictions on the test set
```

```

y_pred = rf_regressor.predict(X_test)

# Evaluate the model
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

# Print the results
print('Mean Squared Error: ',mse)
print('R-squared Score:', r2)

```

Output 🍌

Mean Squared Error: 0.2553684927247781
 R-squared Score: 0.8051230593157366

Process finished with exit code 0

Practical 4F 🍌

```

#Practical 4F - Implement a gradient boosting machine
(e.g., XGBoost).
# Tune hyperparameters and explore feature importance.
from numpy import loadtxt
from xgboost import XGBClassifier
from matplotlib import pyplot
# load data
dataset = loadtxt('pima-indians-diabetes.csv',
delimiter=",")
# split data into X and y
X = dataset[:,0:8]
y = dataset[:,8]
# fit model no training data
model = XGBClassifier()
model.fit(X, y)


```

```
# feature importance
print(model.feature_importances_)
# plot
pyplot.bar(range(len(model.feature_importances_)),
            model.feature_importances_)
pyplot.show()
```

Output :

```
[0.07468373 0.2694981 0.09356754 0.09020481 0.08288945 0.15198407
 0.1017815 0.13539082]
```

Process finished with exit code 0



```
#Practical 5A - Implement and demonstrate the working of  
a Naive Bayesian  
# classifier using a sample data set. Build the model to  
classify a  
# test sample  
  
from sklearn.datasets import load_iris  
from sklearn.model_selection import train_test_split  
from sklearn.naive_bayes import GaussianNB  
from sklearn import metrics  
  
iris = load_iris()  
  
# store the feature matrix (X) and response vector (y)  
X = iris.data  
y = iris.target  
  
# splitting X and y into training and testing sets  
X_train, X_test, y_train, y_test = train_test_split(X, y,  
test_size=0.4,  
random_state=1)  
# training the model on training set  
gnb = GaussianNB()  
gnb.fit(X_train, y_train)  
# making predictions on the testing set  
y_pred = gnb.predict(X_test)  
# comparing actual response values (y_test) with  
predicted response values (y_pred)  
print("Gaussian Naive Bayes model accuracy(in %):",  
      metrics.accuracy_score(y_test, y_pred) * 100)
```

Output :

Gaussian Naive Bayes model accuracy(in %): 95.0

Process finished with exit code 0

```
#Practical 5B- Implement Hidden Markov Models using
hmmlearn

import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from hmmlearn import hmm

# Define the state space
states = ["Silence", "Word1", "Word2", "Word3"]
n_states = len(states)

# Define the observation space
observations = ["Loud", "Soft"]
n_observations = len(observations)

# Define the initial state distribution
start_probability = np.array([0.8, 0.1, 0.1, 0.0])

# Define the state transition probabilities
transition_probability = np.array([[0.7, 0.2, 0.1, 0.0],
                                   [0.0, 0.6, 0.4, 0.0],
                                   [0.0, 0.0, 0.6, 0.4],
                                   [0.0, 0.0, 0.0, 1.0]])

# Define the observation likelihoods
emission_probability = np.array([[0.7, 0.3],
                                  [0.4, 0.6],
```

```
[0.6, 0.4],
[0.3, 0.7]])

# Fit the model
model = hmm.CategoricalHMM(n_components=n_states)
model.startprob_ = start_probability
model.transmat_ = transition_probability
model.emissionprob_ = emission_probability

# Define the sequence of observations
observations_sequence = np.array([0, 1, 0, 0, 1, 1, 0,
1]).reshape(-1, 1)

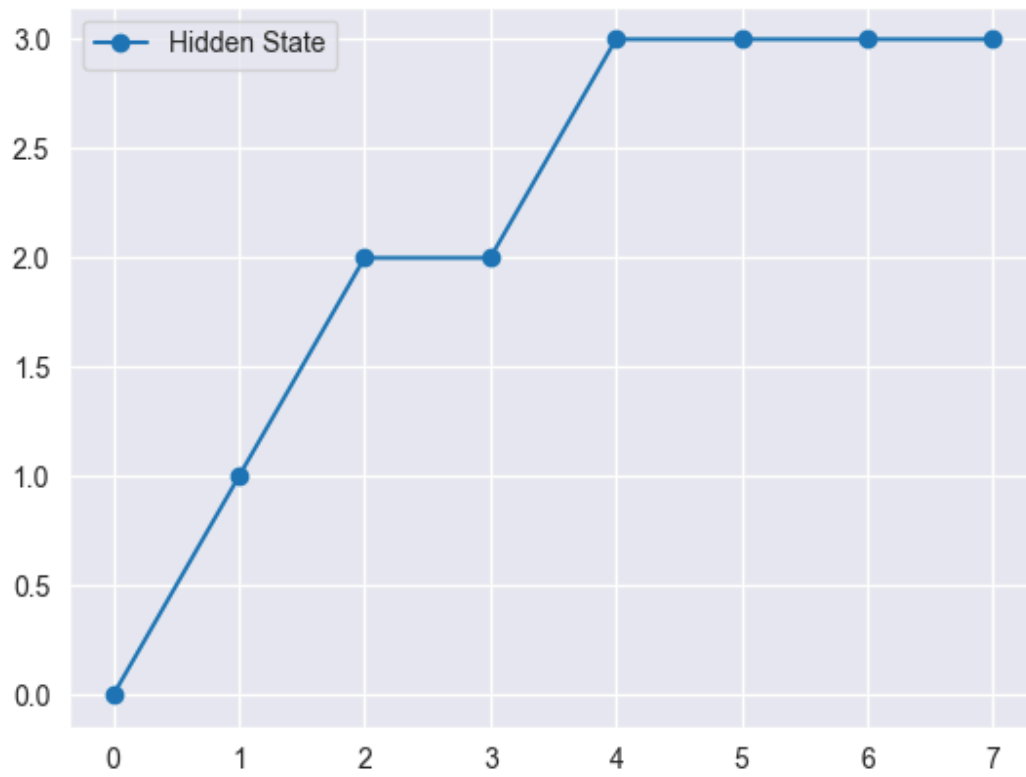
# Predict the most likely hidden states
hidden_states = model.predict(observations_sequence)
print("Most likely hidden states:", hidden_states)

# Plot the results
sns.set_style("darkgrid")
plt.plot(hidden_states, '-o', label="Hidden State")
plt.legend()
plt.show()
```

Output :

Most likely hidden states: [0 1 2 2 3 3 3 3]

Process finished with exit code 0




```
#Practical 6A - Implement Bayesian Linear Regression to  
explore prior  
# and posterior distribution.  
  
from sklearn.datasets import load_iris  
from sklearn.model_selection import train_test_split  
from sklearn.metrics import r2_score  
from sklearn.linear_model import BayesianRidge  
  
# Loading the dataset  
dataset = load_iris()  
  
X = dataset.data  
y = dataset.target  
# Splitting the dataset into testing and training sets  
X_train, X_test, y_train, y_test = train_test_split(X, y,  
                                                    test_size = 0.15,  
                                                    random_state = 42)  
# Creating to train the model  
model = BayesianRidge()  
model.fit(X_train, y_train)  
# Model predicting the test data  
prediction = model.predict(X_test)  
# Evaluation of r2 score of the model against the test  
dataset  
print(f"Test Set r2 score : {r2_score(y_test,  
prediction)}")
```

Output 👍

Test Set r2 score : 0.9225869685714325

Process finished with exit code 0

```
#Practical 6B - Implement Gaussian Mixture Models for  
# density estimation and unsupervised clustering
```

```
import pandas as pd  
import matplotlib.pyplot as plt  
from sklearn import datasets  
from sklearn.mixture import GaussianMixture
```

```
# load the iris dataset  
iris = datasets.load_iris()
```

```
# select first two columns  
X = iris.data[:, :2]
```

```
# turn it into a dataframe  
d = pd.DataFrame(X)
```

```
# plot the data  
plt.scatter(d[0], d[1])  
plt.show()
```

```
gmm = GaussianMixture(n_components = 3)  
gmm.fit(d)
```

```
# Assign a label to each sample  
labels = gmm.predict(d)  
d['labels'] = labels  
d0 = d[d['labels'] == 0]  
d1 = d[d['labels'] == 1]  
d2 = d[d['labels'] == 2]
```

```
# plot three clusters in same plot  
plt.scatter(d0[0], d0[1], c='r')  
plt.scatter(d1[0], d1[1], c='yellow')  
plt.scatter(d2[0], d2[1], c='g')  
plt.show()
```


```
# print the converged log-likelihood value
print('Lower Bound Value is : ',gmm.lower_bound_)
print('Iterations Value is :',gmm.n_iter_)
```


Output 🍌

Lower Bound Value is : -1.4987505566235166

Iterations Value is : 8

Process finished with exit code 0





```

#Practical 7A - Implement cross-validation techniques
# (k-fold, stratified, etc.) for robust model evaluation
# scikit-learn k-fold cross-validation
from numpy import array
from sklearn.model_selection import KFold
# data sample
data = array([0.1, 0.2, 0.3, 0.4, 0.5, 0.6])
# prepare cross validation
kfold = KFold(n_splits=3, shuffle=True, random_state=1)
# enumerate splits
for train, test in kfold.split(data):
    print('train: %s, test: %s' % (data[train],
data[test]))

```

Output 👍

```

train: [0.1 0.4 0.5 0.6], test: [0.2 0.3]
train: [0.2 0.3 0.4 0.6], test: [0.1 0.5]
train: [0.1 0.2 0.3 0.5], test: [0.4 0.6]

```

Process finished with exit code 0

```

# Practical 7B - Systematically explore combinations of
# hyperparameters to optimize model
# performance. (use grid and randomized search)

from sklearn.datasets import load_diabetes
from sklearn.model_selection import train_test_split,
GridSearchCV
from sklearn.neighbors import KNeighborsRegressor
from sklearn.metrics import mean_squared_error

```

```
# Loading the dataset
data = load_diabetes()
X, y = data.data, data.target

# Splitting the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.3, random_state=42)

# Default model
knn_default = KNeighborsRegressor()
knn_default.fit(X_train, y_train)
y_pred_default = knn_default.predict(X_test)
mse_default = mean_squared_error(y_test, y_pred_default)
print(f"Mean Squared Error without tuning:
{mse_default}")

#space to explore
param_grid = {
    'n_neighbors': [3, 5, 7, 9],
    'weights': ['uniform', 'distance'],
    'metric': ['euclidean', 'manhattan']
}
grid_search = GridSearchCV(KNeighborsRegressor(),
param_grid, cv=5, scoring='neg_mean_squared_error')
grid_search.fit(X_train, y_train)
best_knn_grid = grid_search.best_estimator_
y_pred_grid = best_knn_grid.predict(X_test)
mse_grid = mean_squared_error(y_test, y_pred_grid)
print(f"Mean Squared Error with Grid Search tuning:
{mse_grid}")
print(f"Best hyperparameters (Grid Search):
{grid_search.best_params_}")
```

Output 👍

Mean Squared Error without tuning: 3222.117894736842

Mean Squared Error with Grid Search tuning: 3133.022563447985

Best hyperparameters (Grid Search): {'metric': 'euclidean', 'n_neighbors': 9, 'weights': 'distance'}

Process finished with exit code 0

```
#Practical 8 -Implement Bayesian Learning using inferences
```

```
import numpy as np
import matplotlib.pyplot as plt
```

```
# Generate some synthetic data
np.random.seed(42)
true_mu = 5
true_sigma = 2
data = np.random.normal(true_mu, true_sigma, size=100)
```

```
# Define the prior hyperparameters
prior_mu_mean = 0
prior_mu_precision = 1 # Variance = 1 / precision
prior_sigma_alpha = 2
prior_sigma_beta = 2 # Beta = alpha / beta
```

```
# Update the prior hyperparameters with the data
posterior_mu_precision = prior_mu_precision + len(data) /
true_sigma ** 2
posterior_mu_mean=(prior_mu_precision*prior_mu_mean+np.sum(data))/posterior_mu_precision
```

```
posterior_sigma_alpha = prior_sigma_alpha + len(data) / 2
```

```
posterior_sigma_beta = prior_sigma_beta + np.sum((data -
np.mean(data)) ** 2) / 2
```

```
# Calculate the posterior parameters
```

```
posterior_mu = np.random.normal(posterior_mu_mean, 1 /
np.sqrt(posterior_mu_precision), size=10000)
posterior_sigma = np.random.gamma(posterior_sigma_alpha,
1 / posterior_sigma_beta, size=10000)
```

```
# Plot the posterior distributions
```

```
plt.figure(figsize=(10, 4))
plt.subplot(1, 2, 1)
plt.hist(posterior_mu, bins=30, density=True,
color='skyblue', edgecolor='black')
plt.title('Posterior distribution of  $\mu$ ')
plt.xlabel('$\mu$')
plt.ylabel('Density')
```

```
plt.subplot(1, 2, 2)
plt.hist(posterior_sigma, bins=30, density=True,
color='lightgreen', edgecolor='black')
plt.title('Posterior distribution of  $\sigma$ ')
plt.xlabel('$\sigma$')
plt.ylabel('Density')
```

```
plt.tight_layout()
plt.show()
```

```
# Calculate summary statistics
```

```
mean_mu = np.mean(posterior_mu)
std_mu = np.std(posterior_mu)
print("Mean of mu:", mean_mu)
print("Standard deviation of mu:", std_mu)
```

```
mean_sigma = np.mean(posterior_sigma)
std_sigma = np.std(posterior_sigma)
```

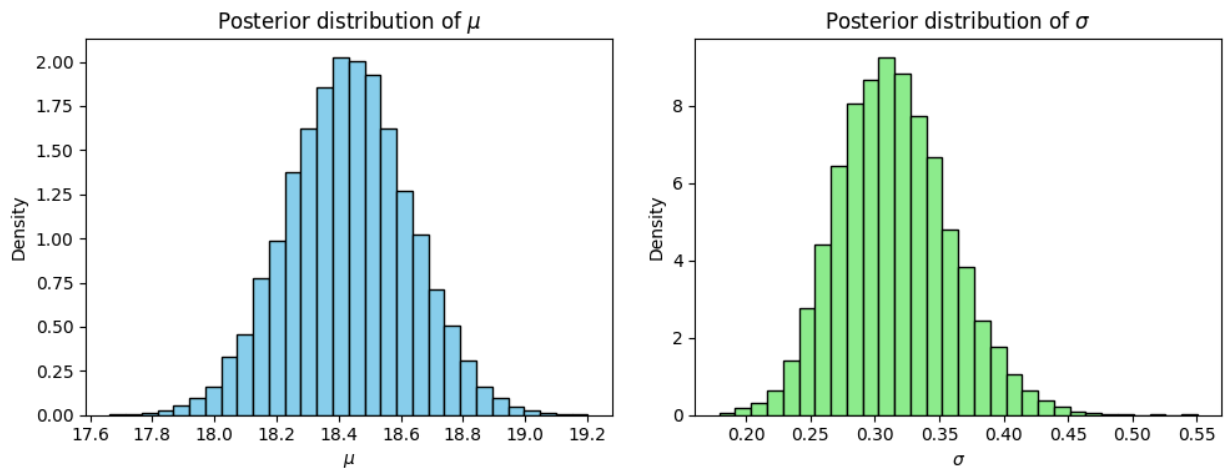


```
print("Mean of sigma:", mean_sigma)
print("Standard deviation of sigma:", std_sigma)
```

Output 🍌

Mean of mu: 18.431718658085725
Standard deviation of mu: 0.19684004415211634
Mean of sigma: 0.314978331043776
Standard deviation of sigma: 0.043338268607342946

Process finished with exit code 0



#Practical 9 - Set up a generator network to produce samples and a discriminator network to

3distinguish between real and generated data. (Use a simple small dataset)

```
import numpy as np
import tensorflow as tf
from tensorflow.keras import layers, Sequential
import matplotlib.pyplot as plt
```

```
def generate_real_data(n_samples):
    x = np.linspace(-1, 1, n_samples)
    y = np.sin(2 * np.pi * x)
    return x, y
```

```
n_samples = 100
x_real, y_real = generate_real_data(n_samples)
```

```
# Visualize the data
plt.plot(x_real, y_real, label="Real Data")
plt.legend()
plt.show()
```

```
def build_generator():
    model = Sequential([
        layers.Dense(16, activation='relu', input_dim=10),
        layers.Dense(16, activation='relu'),
        layers.Dense(2, activation='linear') # Output 2D
    ])
    return model
```

```
generator = build_generator()
generator.summary()
```

```
def build_discriminator():
    model = Sequential([
        layers.Dense(16, activation='relu', input_dim=2),
        layers.Dense(16, activation='relu'),
        layers.Dense(1, activation='sigmoid') # Output
        probability
    ])
    return model

discriminator = build_discriminator()
discriminator.summary()

discriminator.compile(optimizer=tf.keras.optimizers.Adam(
    learning_rate=0.001),
                    loss='binary_crossentropy',
                    metrics=['accuracy'])
```

Output 👍

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 16)	176
dense_1 (Dense)	(None, 16)	272

dense_2 (Dense)	(None, 2)	34
-----------------	-----------	----

Total params: 482 (1.88 KB)

Trainable params: 482 (1.88 KB)

Non-trainable params: 0 (0.00 B)

Model: "sequential_1"


Layer (type)	Output Shape	Param #
dense_3 (Dense)	(None, 16)	48
dense_4 (Dense)	(None, 16)	272
dense_5 (Dense)	(None, 1)	17

Total params: 337 (1.32 KB)

Trainable params: 337 (1.32 KB)

Non-trainable params: 0 (0.00 B)

Process finished with exit code 0



#Practical 10 - Develop an API to deploy your model and perform predictions

Importing Necessary modules

```
from fastapi import FastAPI
import uvicorn
```

Declaring our FastAPI instance

```
app = FastAPI()
```

Defining path operation for root endpoint

```
@app.get('/')
def main():
    return {'message': 'Welcome!'}
```


Defining path operation for /name endpoint

```
@app.get('/{name}')
def hello_name(name : str):
    # Defining a function that takes only string as input
    # and output the
    # following message.
    return {'message': f'Welcome!, {name}'}
```

```
if __name__ == '__main__':
    uvicorn.run(app, host='127.0.0.1', port=8000)
```

```
INFO: Started server process [2904]
INFO: Waiting for application startup.
INFO: Application startup complete.
INFO: Uvicorn running on http://127.0.0.1:8000 (Press CTRL+C to quit)
INFO: 127.0.0.1:62754 - "GET / HTTP/1.1" 200 OK
```

Machine Learning



INFO: Shutting down
INFO: Waiting for application shutdown.
INFO: Application shutdown complete.
INFO: Finished server process [1580]

Process finished with exit code 0